

RESEARCH ARTICLE

Deep Feature Extraction Method for Automatic Classification and Processing of Accounting Information

FENG RUI LIU 

Liaoning Dongding Certified Public Accountants, Benxi 117000, China

e-mail: terkinmeosh@hotmail.com

ABSTRACT In the context of advancing intelligent accounting systems within the broader field of computational methods and data processing, this study focuses on overcoming critical limitations inherent in conventional accounting information classification methodologies. These traditional systems often exhibit low adaptability to rapidly evolving data characteristics and a heavy dependence on manually crafted feature engineering pipelines. Existing approaches predominantly rely on predefined statistical heuristics or shallow machine learning models, which are insufficient for capturing the intricate, non-linear, and high-dimensional relationships embedded in modern financial datasets—especially those generated by large-scale enterprise resource planning systems. To address these constraints, we propose a novel deep feature extraction framework grounded in convolutional autoencoder architectures. This framework is designed to autonomously learn multi-level hierarchical feature representations directly from raw journal entries, eliminating the need for domain-specific feature selection. Our approach incorporates a dual-objective design that simultaneously performs input reconstruction and classification, allowing the network to preserve input fidelity while promoting discriminative learning. Furthermore, an adversarial training component is introduced to enhance generalization performance, especially under conditions of class imbalance and input noise—both common issues in financial transaction data. We evaluate our method on a comprehensive real-world accounting dataset composed of annotated journal logs across multiple business units. Experimental results demonstrate substantial improvements in both classification accuracy and anomaly detection F1-score when compared to traditional baseline models such as logistic regression and decision trees. Our findings illustrate the potential of deep learning-based architectures for intelligent accounting systems and contribute to the growing body of research on adaptive, scalable, and automated solutions for financial data analysis and auditing in computational finance.

INDEX TERMS Deep feature extraction, automatic accounting classification, convolutional autoencoder, anomaly detection, adversarial representation learning.

I. INTRODUCTION

The rapid expansion of enterprise data and the increasing complexity of financial environments have made the automatic classification and processing of accounting information a critical research task. Traditional manual methods are not only labor-intensive but also prone to errors, which can severely impact decision-making quality. As financial

data volume and heterogeneity continue to grow, reliance on manual inspection and rule-based protocols becomes increasingly unsustainable, particularly in large-scale and time-sensitive business operations [1]. Therefore, developing efficient, accurate, and scalable computational methods for handling accounting data is of vital importance. Not only can automated systems reduce operational costs, but they also improve consistency and accuracy in financial reporting by minimizing human error and subjectivity. Automation enables faster reconciliation of accounts, real-time anomaly

The associate editor coordinating the review of this manuscript and approving it for publication was Olarik Surinta .

detection, and compliance monitoring, all of which are essential in today's high-stakes financial ecosystem. Furthermore, with the digitization of financial transactions and the advent of intelligent financial systems, such as enterprise resource planning (ERP) platforms and robotic process automation (RPA), there is a growing demand for advanced techniques that can understand and process both unstructured and structured accounting data in real time [1], [2], [3]. These include journal entries, receipts, invoices, emails, spreadsheets, and transaction logs, which often exhibit irregular formats, semantic ambiguity, and inconsistent quality. This diversity introduces new technical challenges in data representation and processing that traditional systems fail to accommodate. As a result, recent research has increasingly focused on learning-based feature extraction methods that automatically uncover latent structure from complex accounting records without relying on predefined templates. Among these, deep learning-based techniques, particularly those employing hierarchical neural networks, have shown remarkable success in modeling the intricate dependencies present in financial data. This underscores the necessity of exploring and enhancing deep feature extraction methods, which promise a higher level of automation, adaptability, and intelligence in financial data handling, ultimately driving more informed decision-making and resilient financial infrastructures [4], [5].

Initial efforts to address these challenges introduced rule-oriented frameworks that formalized domain expertise into structured processing routines. These systems were built on deterministic rules and heuristic decision trees to emulate accounting logic and label transactions accordingly. By encoding predefined mappings between input patterns and accounting labels, such frameworks offered a degree of transparency and interpretability, which made them suitable for early-stage automation in financial institutions. Their close alignment with established accounting principles also provided a sense of regulatory compliance and traceability, both of which are crucial in the auditing and reporting process. However, these advantages came at the cost of flexibility and scalability [6], [7]. As financial records diversified in format and content—ranging from digital receipts and structured ledger entries to semi-structured emails and invoices—these static systems struggled to maintain performance. The hard-coded rules failed to accommodate edge cases, inconsistencies, and emerging transaction types that deviated from their predefined templates. Moreover, the maintenance of such systems required continuous involvement of domain experts, who needed to update rules and logic in response to policy changes, system migrations, and market innovations. This led to a bottleneck in scalability, as the cost and complexity of upkeep escalated with the volume and variety of data. In dynamic enterprise settings, where thousands of transactions are processed daily across multiple departments and business units, this approach quickly became untenable [8]. The rigidity of rule-based systems also

limited their ability to generalize or learn from historical patterns, rendering them ineffective in scenarios that required contextual awareness or adaptive decision-making. As a result, although rule-oriented frameworks laid an important foundation for automating accounting classification, they highlighted the need for more intelligent systems capable of evolving with data and minimizing dependence on manual rule curation. These limitations ultimately catalyzed the exploration of more flexible computational strategies that could adapt to new data distributions and capture subtle patterns beyond the reach of hand-coded logic [9], [10].

To improve adaptability and reduce dependence on rigid rule structures, statistical learning models were developed to identify classification patterns from annotated financial logs. By utilizing supervised algorithms that learned decision boundaries from examples, these models achieved better generalization across various scenarios, especially when the transaction data exhibited variability in structure and semantics. Techniques such as logistic regression, decision trees, support vector machines, and ensemble methods like random forests became widely adopted in accounting information classification tasks [11], [12]. These approaches offered notable improvements over rule-based systems by allowing models to infer data-label relationships directly from historical observations, rather than relying solely on expert-crafted rules. For instance, models could learn to associate specific account codes with textual descriptions or identify fraud-prone behavior from numerical attributes. This shift introduced a more data-centric paradigm that was faster to deploy across similar task domains and moderately resilient to noise or inconsistencies in input. Nevertheless, despite these advancements, statistical learning models still required significant human intervention in the form of handcrafted features. Designing effective input variables—such as transaction frequency, monetary thresholds, counterparty identifiers, categorical tags, and semantic patterns—demanded deep domain knowledge and iterative refinement. These features often captured only surface-level relationships and failed to reflect deeper contextual interactions among financial entities or temporal sequences [2], [13]. As financial datasets became more complex, encompassing multimodal records like textual justifications, invoice images, and temporal sequences of interrelated entries, conventional feature engineering approaches struggled to keep pace. Moreover, the fixed feature representations limited model flexibility when applied to new institutions or accounting standards, often necessitating manual redesign and retraining. Models also became vulnerable to changes in data distribution or previously unseen transaction types, resulting in performance degradation under real-world deployment conditions. The curse of dimensionality posed further challenges as the number of input features increased. Redundant, noisy, or irrelevant features introduced variance and instability, requiring careful feature selection and regularization techniques to preserve model accuracy [14], [15]. These practical

bottlenecks underscored the necessity of transitioning toward more autonomous and expressive learning frameworks. The increasing demand for models that could ingest raw or minimally processed data while automatically discovering meaningful abstractions set the stage for deeper architectures capable of end-to-end learning without exhaustive manual preprocessing.

In response, a new generation of models has emerged, capable of directly processing raw accounting data through hierarchical representation learning. These models leverage the expressive power of deep neural networks to automatically capture latent structures in complex financial records without the need for exhaustive manual feature engineering. Architectures such as convolutional neural networks (CNNs), recurrent neural networks (RNNs), and transformer-based models have proven particularly effective in this domain. CNNs have been applied to extract spatial correlations in structured accounting matrices or numerical embeddings, while RNNs and gated variants like LSTM and GRU have demonstrated success in modeling the sequential nature of financial transactions, where the order and context of entries often hold critical information. Attention mechanisms have enabled models to selectively focus on salient aspects of the input, facilitating a more nuanced understanding of interrelated entries across long sequences or multi-document records [5], [16]. These advancements have enabled end-to-end learning systems that can process raw journal entries, parse invoice metadata, and even integrate textual annotations or receipts. By learning both local and global patterns, these models excel at tasks such as automatic account classification, fraud detection, and anomaly recognition, achieving higher accuracy and generalization compared to earlier statistical techniques. Furthermore, the use of pretraining on large-scale financial corpora or transfer learning from general language models has enhanced performance in low-resource scenarios, allowing models to adapt to varied accounting vocabularies and practices across organizations [7], [17]. Despite these promising developments, key challenges remain. Deep models often behave as BB, making it difficult for auditors and financial professionals to understand the reasoning behind their outputs. This lack of transparency poses significant concerns in high-stakes settings that demand traceability and compliance with financial regulations. Moreover, the high data requirements and computational costs associated with training such models may limit their accessibility for smaller enterprises or real-time applications. Consequently, ongoing research is increasingly focused on improving interpretability through techniques such as saliency mapping, attribution scores, and prototype learning, as well as enhancing data efficiency via semi-supervised learning, knowledge distillation, and compact network design. These efforts aim to bridge the gap between high model performance and practical usability in modern accounting intelligence systems [10], [18].

To address the limitations of previous approaches—including symbolic systems' lack of adaptability, machine learning's dependence on manual features, and deep models' computational overhead—we propose a refined feature extraction framework tailored to the automatic classification and processing of accounting information. This method integrates domain-specific attention mechanisms with efficient deep architectures to capture both global transaction trends and local entry nuances. Unlike prior models, which either focused narrowly on structured data or required extensive labeled corpora, our approach balances generalization with specificity, adapting across diverse accounting contexts with minimal supervision. Moreover, it incorporates lightweight transfer learning strategies, enabling effective performance with limited computational resources. This shift toward more specialized yet scalable solutions is essential for meeting the increasing demands of intelligent financial information systems.

In summary, our proposed framework offers a robust and intelligent solution to the challenges in automated accounting information processing. Below, we summarize the **main innovations** of this work:

- **Domain-Attentive Graph Modeling:** We introduce a novel graph encoder that captures temporal evolution, semantic structure, and transaction dependencies across financial records.
- **Audit-Informed Reinforcement Learning:** Our AIRP module integrates compliance constraints directly into the policy learning process, enabling interpretable and regulation-aware fiscal decision-making.
- **Hybrid Symbolic-Neural Integration:** The model connects symbolic audit rules with deep learning architectures, bridging the gap between transparency and predictive performance.

These innovations collectively address the limitations of rule-based and shallow models, establishing a new direction for scalable, explainable, and compliant accounting automation.

II. RELATED WORK

A. DEEP LEARNING IN ACCOUNTING

Recent advancements in deep learning have significantly transformed the landscape of accounting information systems by enabling automatic classification and intelligent processing of complex financial data. One of the most influential applications has been the utilization of neural networks, particularly convolutional neural networks (CNNs) and recurrent neural networks (RNNs), to extract semantic and structural features from large volumes of unstructured or semi-structured accounting data. These models can process various forms of financial inputs, such as invoices, ledgers, balance sheets, and narrative reports, identifying patterns that are imperceptible to traditional statistical methods [1], [19]. The integration of word embeddings like Word2Vec, GloVe, and BERT within these frameworks has allowed models to

capture contextual and domain-specific meanings of accounting terms, enhancing classification accuracy. Studies such as those by Zhang et al. [38] have shown that deep learning outperforms rule-based and machine learning approaches in fraud detection and audit risk assessment. Similarly, autoencoders have been applied to compress high-dimensional accounting features into compact representations suitable for downstream classification tasks. These representations often capture latent variables such as transaction regularities, cost structures, or anomaly indicators that are critical for financial decision-making [12], [20]. Furthermore, recent attention-based models and architectures, such as BERT and its accounting-specific adaptations, have demonstrated superior capabilities in handling long-range dependencies in financial narratives, allowing for more nuanced document classification and sentiment analysis. Integration with optical character recognition (OCR) tools has extended these capabilities to scanned and handwritten financial documents, providing end-to-end automation pipelines [13], [21]. However, challenges remain in ensuring interpretability, regulatory compliance, and robustness to adversarial inputs, particularly in high-stakes accounting environments. The field continues to evolve towards hybrid models that combine rule-based domain knowledge with deep representations, aiming to preserve explainability while leveraging the power of neural computation [15], [22].

Recent advances in the application of deep learning and explainable AI (XAI) have significantly influenced the accounting and auditing domains. Zhang et al. emphasize the growing importance of explainability in auditing processes, highlighting how XAI techniques can improve trust, transparency, and regulatory alignment in financial systems [23]. Similarly, Wu and Du demonstrate that deep learning models can effectively detect financial statement fraud in complex and data-rich environments, such as those found in Chinese listed firms [24]. Craja et al. further explore various deep learning architectures and show their superiority over traditional statistical methods in identifying fraudulent patterns in financial disclosures [25]. These studies underscore the potential of AI in enhancing accounting practices. Our work builds on these foundations by incorporating constraint-aware learning and temporal-structural graph modeling to support real-time and regulation-sensitive financial decision-making.

B. FINANCIAL DOCUMENT UNDERSTANDING

The domain of financial document understanding has seen remarkable progress with the advent of deep feature extraction techniques tailored to the unique structure and semantics of accounting texts. Unlike general-purpose document classification tasks, accounting documents often contain heterogeneous data types including numerical tables, textual narratives, and structured metadata. This complexity necessitates specialized architectures that can jointly process multimodal inputs [16], [26]. State-of-the-art systems have

employed hierarchical attention mechanisms to model the interdependence between document sections such as headers, line items, footnotes, and annotations. For instance, methods like LayoutLM and its successors integrate visual layout information with text embeddings to better understand document context, which is particularly crucial in financial statements where spatial positioning conveys semantic roles. Graph neural networks (GNNs) have also emerged as a powerful tool for modeling relationships between accounting entities, enabling the discovery of latent links such as related party transactions or networked financial dependencies [17], [27]. In tasks such as automatic journal entry classification, entity recognition, and audit trail reconstruction, deep features extracted from pre-trained language models are fine-tuned to capture domain-specific regularities. Moreover, cross-document attention and retrieval models allow systems to contextualize entries by referencing related documents, such as linking an expense claim to a contract or procurement record. Despite these advances, generalization remains a significant hurdle due to the diversity in document formats, accounting standards, and terminologies across organizations and jurisdictions [28], [29]. Efforts are underway to construct large-scale, annotated corpora of accounting documents and to develop domain-adaptive pretraining methods that align linguistic and numerical patterns with accounting logic. The pursuit of explainable AI (XAI) in financial document analysis is also gaining momentum, with attention visualization and concept attribution being employed to enhance model transparency and trustworthiness for auditors and regulators [1], [30].

C. AUTOMATED ACCOUNTING CLASSIFICATION

Automated classification of accounting entries and transactions is a foundational component of intelligent financial systems, and recent deep learning techniques have markedly improved performance in this domain by enabling sophisticated feature extraction from raw and semi-structured inputs. Traditional accounting systems rely heavily on manually crafted rules and extensive domain knowledge, which are both labor-intensive and inflexible to evolving data landscapes [19]. In contrast, modern systems utilize end-to-end deep learning pipelines that learn hierarchical representations of transactions directly from data, capturing intricate dependencies between attributes such as date, amount, category, vendor, and narrative description. Transformer-based models, particularly those adapted for tabular and sequence data, have demonstrated exceptional capabilities in capturing both local syntactic features and global semantic patterns in accounting records. Sequence-to-sequence models have been applied to tasks such as account mapping, VAT code classification, and regulatory tagging, while attention mechanisms ensure that critical transaction features receive higher importance during classification [20], [31]. Transfer learning and domain adaptation techniques further enhance these models by incorporating financial

knowledge from publicly available datasets and adapting to organization-specific idiosyncrasies [21]. Integration with external knowledge bases, such as accounting standards or taxonomies, allows for constraint-aware classification that respects regulatory frameworks. Reinforcement learning and active learning paradigms have also been proposed to iteratively refine classification strategies based on feedback from human accountants, improving model performance over time. Despite these advancements, challenges persist in handling noisy data, imbalanced class distributions, and ambiguous transaction descriptions [22]. Research is also focusing on continual learning approaches that allow classification systems to update their models in response to new data without catastrophic forgetting. The convergence of deep feature extraction, domain-specific knowledge integration, and interactive learning is poised to redefine the automation landscape of accounting information processing, paving the way for more efficient, accurate, and adaptive financial systems [26].

III. METHOD

A. OVERVIEW

To clearly articulate the innovative contributions of this work, we emphasize three core aspects that distinguish our framework from prior methods. Our model introduces a novel domain-attentive dynamic graph encoder that captures not only the temporal evolution of financial transactions but also the semantic and hierarchical roles of accounting entities, enabling fine-grained feature extraction beyond traditional sequence-based models. We embed regulatory and compliance constraints directly into the learning and decision-making pipeline through a constraint-aware optimization layer and audit-guided reinforcement planning (AIRP), which ensures both accuracy and accountability in fiscal actions. Our hybrid integration of symbolic audit rules with deep neural representations bridges the gap between explainability and performance, addressing a long-standing challenge in financial automation. These innovations jointly contribute to a scalable, interpretable, and regulation-aware framework for intelligent accounting information processing.

Enterprise financial management (EFM) stands at the core of organizational sustainability and strategic progression, integrating a wide array of fiscal activities to ensure operational stability, regulatory compliance, and long-term value creation. As businesses navigate increasingly volatile global markets, the demand for intelligent, data-driven financial systems has intensified, prompting a shift from traditional accounting practices to comprehensive, digitized financial ecosystems. This paper presents a structured methodological framework to advance the analytical depth and operational precision of enterprise financial management through a formal, model-driven, and algorithmic lens.

The forthcoming sections delineate the foundational constructs, architectural enhancements, and strategic mechanisms proposed in this work. In Section III-B, we establish the formal underpinnings of enterprise financial processes,

defining a unified notation system and mathematical formulation to abstract and model the dynamics of financial assets, liabilities, operational flows, and compliance constraints. This section lays the groundwork by expressing key financial interactions as structured operators over time-indexed variable spaces, enabling symbolic reasoning and analytical tractability. In Section III-C, we introduce our novel architectural formulation, titled FinGraphNet, which reconceptualizes financial management as a graph-based temporal reasoning task. FinGraphNet encodes multi-source financial data—including budgetary forecasts, real-time transaction streams, and compliance logs—into an adaptive, structured graph representation. By coupling dynamic embedding modules with constrained optimization layers, our model supports both fine-grained inference and high-level fiscal decision modeling. This structural model diverges from conventional ledger-based systems by emphasizing temporal correlation, dependency propagation, and structural regularization across fiscal entities. Section III-D articulates our strategic engine, coined Audit-Informed Reinforcement Planning (AIRP), designed to leverage domain-specific audit trails and control signals to steer financial decision paths. AIRP constructs a knowledge-driven feedback mechanism into the management loop by extracting implicit constraints from historical audit records and encoding them into reinforcement learning agents. These agents simulate counterfactual fiscal scenarios and optimize decision policies under uncertainty and compliance constraints, making the financial system not only reactive but also anticipatory. Collectively, these three components construct a cohesive methodology that transcends static financial reporting, aligning financial intelligence with organizational agility and regulatory robustness. By embedding semantic structure, algorithmic interpretability, and strategic feedback into enterprise financial operations, this approach proposes a scalable and formally grounded pathway to modernize EFM practices. The remainder of this section outlines the critical roles and interconnections among the upcoming subsections, offering a roadmap for the methodological innovations that follow.

To ensure clarity for readers from both technical and accounting backgrounds, we include a glossary of key financial terms used throughout the paper. These definitions provide concise explanations of fundamental concepts—such as assets, liabilities, liquidity, and audit constraints—that are central to our modeling framework. The glossary (see Table 1) aims to assist interdisciplinary readers in understanding the domain-specific context in which the proposed methods operate.

B. PRELIMINARIES

Enterprise Financial Management (EFM) can be formalized as a multi-layered dynamic system operating over discrete fiscal periods. Let $\mathcal{T} = \{t_1, t_2, \dots, t_N\}$ denote a sequence of time-indexed decision epochs, where each t_i represents the closure of a financial reporting cycle. At the core of our

TABLE 1. Glossary of key accounting terms.

Term	Description
Asset (A_t)	Resources owned by a company that are expected to generate future economic value.
Liability (L_t)	Financial obligations the company must settle in the future (e.g., loans, accounts payable).
Revenue (R_t)	Income generated from core business operations during a specific time period.
Cost (C_t)	Expenses incurred to generate revenue, including operating and capital expenditures.
Equity (E_t)	Residual interest in assets after deducting liabilities; represents ownership value.
Liquidity (LQ_t)	A measure of a firm's ability to meet its short-term obligations using available assets.
Debt Ratio	The proportion of total assets financed through debt, used to assess financial leverage.
Coverage Threshold	A ratio ensuring earnings are sufficient to cover debt payments and interest obligations.
Audit Constraint	Regulatory or policy-based rule derived from audits, restricting financial behavior.
Compliance Bound	Boundary condition ensuring financial actions remain within legal or policy frameworks.

formulation is the modeling of an enterprise's financial state as a structured tuple over a time-varying domain.

We define the financial state at time t as,

$$\mathcal{S}_t = (A_t, L_t, R_t, C_t, E_t), \quad (1)$$

where A_t denotes the asset vector, L_t the liability vector, R_t the revenue flow, C_t the cost flow, and E_t the equity position at time t .

To reflect financial operations, we define the net cash transformation operator and liquidity evolution as,

$$LQ_{t+1} = LQ_t + \left(\sum_{i=1}^n R_t^{(i)} - \sum_{j=1}^m C_t^{(j)} \right) - \mathcal{D}_t, \quad (2)$$

where $\mathcal{D}_t$ represents debt servicing obligations.

The intertemporal objective of EFM integrates discounted utility of financial strategies,

$$\mathcal{J} = \mathbb{E} \left[\sum_{t=1}^T \gamma^t U(\mathcal{S}_t, \mathcal{U}_t) \right], \quad (3)$$

where U is a utility function, and $\gamma \in (0, 1)$ is the discount factor.

We impose regulatory compliance via constraint functions,

$$\mathcal{C}_t^{(\ell)} : g_\ell(\mathcal{S}_t, \mathcal{U}_t) \leq 0, \quad \ell = 1, \dots, L, \quad (4)$$

capturing rules on debt ratios, coverage thresholds, and exposure limits.

The planning objective is to determine optimal decisions $\{\mathcal{U}_t\}_{t=1}^T$ over a feasible set,

$$\max_{\{\mathcal{U}_t\}} \mathcal{J} \quad \text{subject to} \quad \mathcal{C}_t^{(\ell)} \text{ and fiscal constraints.} \quad (5)$$

C. FIN GRAPH NET

We introduce FinGraphNet, a structured graph-based model tailored to enterprise financial management, which integrates multivariate temporal data, fiscal structure semantics, and regulatory logic into a unified learning and reasoning framework. The model treats financial entities and their interactions as nodes and edges in a dynamic, directed multigraph $\mathcal{G}_t = (\mathcal{V}_t, \mathcal{E}_t)$ at each time step $t \in \mathcal{T}$ (As shown in Figure 1).




FIGURE 1. High-level architecture of FinGraphNet. The process begins with raw accounting data, which is transformed into a time-varying financial graph structure. Dynamic graph encoding computes contextualized node and edge representations (Equations 3–7). These representations are propagated through a temporal-aware message passing mechanism with attention (Equations 8–10), followed by financial memory updates via recurrent and smoothing layers (Equations 11–13). A constraint-aware projection layer ensures regulatory feasibility (Equations 14–16), and the output scores are used to guide reinforcement planning (Equations 17–18).

1) DYNAMIC GRAPH ENCODING

Let $\mathcal{V}_t = \{v_1^{(t)}, v_2^{(t)}, \dots, v_n^{(t)}\}$ denote the set of financial entities at time t , where each node represents a financial item (As shown in Figure 2).

Each node $v_i^{(t)}$ is associated with a multivariate feature vector $x_i^{(t)} \in \mathbb{R}^d$, constructed from both temporal and structural attributes,

$$x_i^{(t)} = \phi(v_i^{(t)}) = [\text{hist}_i^{(t)}, \text{sector}_i, \text{risk}_i, \text{flow}_i^{(t)}], \quad (6)$$

where $\text{hist}_i^{(t)}$ captures trailing time-series indicators, sector_i encodes the hierarchical role of the node, risk_i represents

Comparative Study: Performance Metrics on Orbis and CSMAR Datasets




FIGURE 2. Detailed view of the dynamic graph encoder. This figure shows how raw financial entities are transformed into graph representations. Node embeddings are constructed using historical trends, sector hierarchy, risk weights, and cash flows (Equation 3). Edge features encode directional relationships using transformation matrices, semantic relations, and elasticity (Equation 5). Bilateral edge constraints (Equation 7) maintain consistency in pairwise relations.

systemic exposure weights, and $\text{flow}_i^{(t)}$ quantifies net cash movement during the current epoch.

We define the normalized node degree as an auxiliary regularizer,

$$d_i^{(t)} = \frac{1}{|\mathcal{N}_i|} \sum_{j \in \mathcal{N}_i} \|\text{flow}_{ij}^{(t)}\|_1, \quad (7)$$

where $\mathcal{N}_i$ is the neighborhood of $v_i^{(t)}$ and $\text{flow}_{ij}^{(t)}$ represents directed cash flow magnitude between v_i and v_j .

The edge set $\mathcal{E}_t = \{(v_i^{(t)} \rightarrow v_j^{(t)}, e_{ij}^{(t)})\}$ models directed relationships between financial nodes, capturing credit transfers, investment dependencies, or regulatory transfers. Each edge embedding includes a transformation matrix $W_{ij}^{(t)} \in \mathbb{R}^{d \times d}$ and a semantic relation vector $r_{ij} \in \mathbb{R}^k$,

$$e_{ij}^{(t)} = \psi(v_i^{(t)}, v_j^{(t)}) = [W_{ij}^{(t)}, r_{ij}, \lambda_{ij}^{(t)}], \quad (8)$$

where $\lambda_{ij}^{(t)}$ is a learned scalar denoting the elasticity of the financial interaction under stress scenarios.

To enforce directed consistency, we introduce a bilateral edge constraint,

$$\psi(v_i^{(t)}, v_j^{(t)}) + \psi(v_j^{(t)}, v_i^{(t)}) \approx \mathcal{I}_{ij}^{(t)}, \quad (9)$$

where $\mathcal{I}_{ij}^{(t)}$ denotes an identity-preserving neutral transfer operator, facilitating antisymmetric balance encoding.

The inter-temporal continuity of financial behavior is captured using a decay-regularized update of node features,

$$x_i^{(t)} = \gamma x_i^{(t-1)} + (1 - \gamma) \cdot \phi(v_i^{(t)}), \quad (10)$$

with decay coefficient $\gamma \in (0, 1)$ controlling the memory of historical feature dynamics.

To enhance structural granularity, we generate a composite edge strength metric across relation types,

$$\zeta_{ij}^{(t)} = \sum_{\tau=1}^k \omega_{\tau} \cdot r_{ij}^{(\tau)}, \quad (11)$$

where ω_{τ} are trainable weights for each relation dimension τ .

This dynamic encoding enables the formation of time-varying multigraphs that evolve with fiscal signals, allowing downstream models to access structurally enriched, temporally grounded financial representations.

2) TEMPORAL-AWARE PROPAGATION

FinGraphNet propagates information across financial graphs through a temporally contextualized, relation-aware message passing mechanism that enables both short- and long-range dependency modeling. Each node $v_i^{(t)}$ aggregates messages from its local neighborhood $\mathcal{N}(i)$ using an attention-weighted scheme,

$$h_i^{(t+1)} = \sigma \left(\sum_{j \in \mathcal{N}(i)} \alpha_{ij}^{(t)} W_{ij}^{(t)} h_j^{(t)} + U x_i^{(t)} + b_i \right), \quad (12)$$

where $h_j^{(t)}$ is the previous hidden state of neighbor v_j , $W_{ij}^{(t)}$ is the edge-type-specific transformation matrix, U is a learnable linear projection applied to the current node's features, b_i is a bias term, and σ is a nonlinear activation such as ReLU or GELU.

The attention coefficient $\alpha_{ij}^{(t)}$ reflects the relevance of each neighbor v_j to v_i at time t , jointly considering structural relation embeddings and latent node states,

$$\alpha_{ij}^{(t)} = \frac{\exp \left(a^{\top} \tanh(W_q h_i^{(t)} + W_k h_j^{(t)} + W_r r_{ij} + W_f f_{ij}^{(t)}) \right)}{\sum_{k \in \mathcal{N}(i)} \exp \left(a^{\top} \tanh(W_q h_i^{(t)} + W_k h_k^{(t)} + W_r r_{ik} + W_f f_{ik}^{(t)}) \right)}, \quad (13)$$

where $f_{ij}^{(t)}$ is a flow-contextual edge feature such as recent transaction activity or cumulative volume, and W_f its transformation matrix. This attention integrates semantic and quantitative relation awareness.

To incorporate temporal continuity, FinGraphNet models the evolution of hidden states with a gated recurrent mechanism,

$$s_i^{(t+1)} = \text{GRU}(h_i^{(t+1)}, s_i^{(t)}), \quad (14)$$

where $s_i^{(t)}$ is a time-evolving latent vector that encodes financial memory of account v_i . It captures both latent behavioral drift and transaction periodicity, enabling recurrent aggregation of contextual trends.

We introduce a temporal smoothing mechanism to mitigate abrupt transitions and amplify stable dynamics,

$$\bar{s}_i^{(t+1)} = \beta s_i^{(t+1)} + (1 - \beta) \cdot \bar{s}_i^{(t)}, \quad (15)$$

where $\bar{s}_i^{(t)}$ is the smoothed state vector and $\beta \in (0, 1)$ is a smoothing coefficient. This representation is used for downstream decision modules requiring trend-consistent features.

By layering multiple propagation and recurrent units across epochs, the model builds hierarchical node embeddings that are responsive to evolving fiscal conditions and inter-entity dependencies.

3) CONSTRAINT-AWARE OPTIMIZATION

To ensure that predicted financial flows adhere to regulatory constraints, we formulate the problem as a constrained projection task. Given a predicted flow vector $\tilde{x}^{(t+1)}$ that may violate constraints, our goal is to find the closest feasible flow vector $\hat{x}^{(t+1)}$ within the admissible region $\mathcal{F}$ defined by linear inequalities. This is formally expressed as a quadratic optimization problem. First, we define the feasible region as:

$$\mathcal{F} = \{x \in \mathbb{R}^d \mid Gx \leq h\}, \quad (16)$$

where $G \in \mathbb{R}^{L \times d}$ is a matrix encoding L linear constraints (e.g., budget ceilings, debt coverage ratios), and $h \in \mathbb{R}^L$ is the corresponding threshold vector. Then, to find the closest point within $\mathcal{F}$ to the unconstrained prediction $\tilde{x}^{(t+1)}$, we solve the following optimization problem:

$$\hat{x}^{(t+1)} = \arg \min_{x \in \mathcal{F}} \|x - \tilde{x}^{(t+1)}\|_2^2. \quad (17)$$

This problem minimizes the squared Euclidean distance between the original prediction and the feasible point, subject to linear inequality constraints. It is a standard Quadratic Programming (QP) problem, which is convex and can be solved efficiently using projection methods or QP solvers.

To respect financial regulations and policy constraints, we introduce a constraint-satisfaction projection layer. Given predicted flows $\tilde{x}^{(t+1)}$, we enforce feasibility via,

$$\hat{x}^{(t+1)} = \arg \min_{x \in \mathcal{F}} \|x - \tilde{x}^{(t+1)}\|_2^2, \quad (18)$$

where $\mathcal{F} = \{x \mid Gx \leq h\}$ encodes budget limits, liquidity reserves, and compliance bounds using matrix inequality representations.

To evaluate and prioritize financial strategies, we define a scoring function,

$$\text{Score}(v_i^{(t)}) = \rho^\top s_i^{(t)} - \kappa^\top \xi_i^{(t)}, \quad (19)$$

where ρ and κ are learnable risk/reward weights, and $\xi_i^{(t)}$ is a vector of estimated exposure metrics.

The model is trained on historical transaction data with supervisory signals such as future cash position $\hat{L}Q_{t+\tau}$ and realized value creation $\hat{E}_{t+\tau}$,

$$\mathcal{L}_{\text{reg}} = \sum_i \left\| \hat{L}Q_{t+\tau} - LQ_{t+\tau}^{\text{pred}} \right\|_2^2 + \eta \left\| \hat{E}_{t+\tau} - E_{t+\tau}^{\text{pred}} \right\|_2^2, \quad (20)$$

where η balances liquidity versus equity tracking.

We penalize constraint violations using a hinge-like loss,

$$\mathcal{L}_{\text{cons}} = \sum_{\ell=1}^L \max(0, g_\ell(\hat{x}^{(t+1)}))^2. \quad (21)$$

The full training objective integrates prediction accuracy and constraint adherence,

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{reg}} + \lambda \mathcal{L}_{\text{cons}} + \mu \|\Theta\|_2^2, \quad (22)$$

where Θ denotes all model parameters and λ, μ are regularization hyperparameters.

D. AUDIT-INFORMED REINFORCEMENT PLANNING (AIRP)

To operationalize the predictive insights of FinGraphNet into adaptive fiscal decision-making, we propose a novel strategy module named Audit-Informed Reinforcement Planning (AIRP). This mechanism integrates historical audit data and compliance records into a reinforcement learning (RL) framework, enabling strategic financial control under constraints and uncertainty (As shown in Figure 3).




FIGURE 3. Audit-Informed Reinforcement Planning (AIRP) module. This figure details the AIRP mechanism, where FinGraphNet embeddings are used to generate candidate financial actions. These actions are checked against audit-informed constraints (Equation 24) and projected into a feasible set via quadratic programming (Equation 25). Audit conformity and slack losses (Equations 26–27) are integrated into the final policy loss (Equation 29), enabling compliant and adaptive decision-making.

1) CONSTRAINT-AWARE MDP FORMULATION

We model the decision environment as a constrained Markov Decision Process (MDP) $\mathcal{M} = (\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \Gamma)$, tailored to financial planning under structural, regulatory, and risk-aware constraints. The state space $\mathcal{S}$ consists of latent financial embeddings $s_t \in \mathbb{R}^d$ generated by FinGraphNet, capturing the evolving economic posture of an enterprise. The action space $\mathcal{A}$ contains fiscal actions such as allocation strategies, denoted $\mathcal{A}_t \in \mathbb{R}^m$, which include investment proportions, debt repayments, and liquidity buffers (As shown in Figure 4).

State transitions follow parameterized financial dynamics,

$$s_{t+1} = f_\phi(s_t, \mathcal{A}_t, \omega_t), \quad (23)$$




FIGURE 4. Constraint-Aware MDP Cell within AIRP. This diagram outlines a single decision step in the MDP formulation of AIRP. A candidate action is sampled from the current policy and evaluated through a constraint-aware transition function (Equation 21). The reward is computed by incorporating audit violations and downside risk (Equations 23–27), and the policy is updated using a trust-region method (Equation 29) to ensure stable optimization.

where f_ϕ is a differentiable transition function, and ω_t denotes stochastic exogenous shocks, such as inflation, tax policy shifts, or demand volatility.

The stochastic policy $\pi_\theta(\mathcal{A}_t | s_t)$ maps financial states to action distributions and is optimized to maximize the expected discounted cumulative reward,

$$J(\theta) = \mathbb{E}_{\pi_\theta} \left[\sum_{t=1}^T \gamma^t R(s_t, \mathcal{A}_t) \right], \quad (24)$$

subject to the feasibility constraint $\mathcal{A}_t \in \Gamma_t(s_t)$, where Γ_t denotes the admissible action space derived from historical audit constraints and structural fiscal rules.

The reward function is shaped to reflect multiple enterprise performance indicators and penalize risk and non-compliance,

$$\begin{aligned} R(s_t, \mathcal{A}_t) \\ = v_1 \cdot \text{ROE}_t + v_2 \cdot \text{LCR}_t - v_3 \cdot \text{Risk}_t - v_4 \cdot \text{Violation}_t, \end{aligned} \quad (25)$$

where return on equity (ROE), liquidity coverage ratio (LCR), and constraint violations are computed from both predicted cash flow models and real-time regulatory assessments.

Constraint violation is further formalized as a soft penalty function,

$$\text{Violation}_t = \sum_{\ell=1}^L \max(0, c_\ell^\top \mathcal{A}_t - d_\ell)^2, \quad (26)$$

where (c_ℓ, d_ℓ) encodes the ℓ -th linear inequality derived from audit rules or internal control policies.

To ensure conservative fiscal behavior under uncertainty, we augment the reward with a downside risk term,

$$\text{Risk}_t = \mathbb{E}_{\omega_t} \left[\|s_{t+1} - \mathbb{E}[s_{t+1}]\|_2^2 \right], \quad (27)$$

which captures the second-order moment of the predicted state, penalizing volatile trajectories.

The admissible action set $\Gamma_t(s_t)$ is dynamically extracted from compliance records, ensuring that policy generation remains in alignment with legal and strategic boundaries,

$$\Gamma_t(s_t) = \{\mathcal{A}_t \in \mathbb{R}^m \mid G_t \mathcal{A}_t \leq h_t\}, \quad (28)$$

where (G_t, h_t) represents time-indexed constraint matrices, derived from past audit documents and risk simulations.

2) AUDIT-GUIDED ACTION PROJECTION

To ensure fiscal decision-making adheres to regulatory and historical compliance patterns, we incorporate audit-informed constraints into the policy execution phase. Historical audit logs $\mathcal{A}_{\log}$ are mined using pattern extraction algorithms and domain templates to derive interpretable linear constraints that formalize financial control boundaries. The extracted set of admissible actions at time t is denoted as,

$$\Gamma_t = \left\{ \mathcal{A}_t \in \mathbb{R}^m \mid \forall \ell : c_\ell^\top \mathcal{A}_t \leq d_\ell \right\}, \quad (c_\ell, d_\ell) \in \text{Audit}_t, \quad (29)$$

where each (c_ℓ, d_ℓ) represents a constraint induced from audit triggers such as overspending on capital, excessive leverage, or underfunding of reserves.

Given an action proposal $\mathcal{A}_t$ sampled from the policy π_θ , it is not directly executed but projected back into the feasible region Γ_t to ensure constraint satisfaction,

$$\tilde{\mathcal{A}}_t = \arg \min_{\mathcal{A} \in \Gamma_t} \|\mathcal{A} - \mathcal{A}_t\|_2^2, \quad (30)$$

which is a quadratic program solvable in closed-form or via projected gradient descent, depending on the complexity of Γ_t .

To operationalize soft constraint adherence during training, we define an audit-conformity regularization loss that penalizes deviations from critical hyperplanes,

$$\mathcal{L}_{\text{audit}} = \sum_t \sum_{\ell=1}^{L_t} \max(0, \Delta_{t,\ell}^\top \mathcal{A}_t - b_{t,\ell})^2, \quad (31)$$

where $(\Delta_{t,\ell}, b_{t,\ell})$ denotes the ℓ -th regulatory hyperplane at time t , reflecting either statutory thresholds or internal compliance metrics extracted from flagged entries in $\mathcal{A}_{\log}$.

Furthermore, to allow dynamic flexibility while retaining audit awareness, we define a soft-bound relaxation penalty,

$$\mathcal{L}_{\text{slack}} = \sum_t \sum_{\ell=1}^{L_t} \xi_{t,\ell}^2, \quad \text{s.t.} \quad \Delta_{t,\ell}^\top \mathcal{A}_t \leq b_{t,\ell} + \xi_{t,\ell}, \quad (32)$$

where $\xi_{t,\ell}$ represents slack variables that permit bounded constraint violations for optimization flexibility, and are regularized to remain small.

To unify constraint projection and learning, the adjusted action is incorporated into the actor update via a dual objective,

$$\mathcal{L}_{\text{proj}} = \|\tilde{\mathcal{A}}_t - \mathcal{A}_t\|_2^2 + \alpha \cdot \mathcal{L}_{\text{audit}} + \beta \cdot \mathcal{L}_{\text{slack}}, \quad (33)$$

where α and β are tunable weights balancing hard and soft regulatory conformance. This enables reinforcement strategies to remain proactive while respecting compliance boundaries throughout learning and execution.

3) RISK-DRIVEN POLICY LEARNING

We estimate financial volatility via conditional variance,

$$\text{Risk}_t = \text{Tr}(\mathbb{V}[\mathcal{S}_{t+1} \mid s_t, \mathcal{A}_t]), \quad (34)$$

using empirical bootstrapping over FinGraphNet predictions.

AIRP includes a counterfactual simulator $\mathcal{C}$ estimating alternate fiscal paths,

$$\hat{s}_{t+\tau}^{\text{cf}} = \mathcal{C}(s_t, \mathcal{A}_t^{\text{alt}}, \tau), \quad (35)$$

used to evaluate hypothetical decisions under historical stress scenarios.

We adopt a trust-region method for robust learning,

$$\theta \leftarrow \arg \max_{\theta'} J(\theta') - \eta \cdot \text{KL}(\pi_{\theta'} \parallel \pi_\theta), \quad (36)$$

where η controls the step size in policy space.

The full training loss incorporates strategic utility, risk, and audit alignment,

$$\mathcal{L}_{\text{total}} = -J(\theta) + \alpha \cdot \mathcal{L}_{\text{audit}} + \beta \cdot \mathbb{E}[\text{Risk}_t], \quad (37)$$

where α, β weight regulatory and stability components.

IV. EXPERIMENTAL SETUP

A. DATASET

The Compustat dataset [29] provides comprehensive time-series financial data, emphasizing projected budgets across departments. It facilitates the study of long-term fiscal planning and inter-departmental resource allocation. With quarterly forecasts, historical expenditures, and relevant economic indicators, the dataset captures fluctuations in budget estimations driven by policy shifts or macroeconomic conditions. It is well-suited for tasks such as temporal modeling, budget anomaly detection, and multivariate regression, and includes metadata like region, sector, and reporting authority for enriched contextual analysis. The EDGAR dataset [30] offers daily-level financial transaction data across diverse user profiles and spending categories, making it

particularly valuable for personal finance management and consumer behavior analysis. It encompasses both structured and unstructured information—such as timestamps, vendor names, transaction modes, and categorical labels—enabling a wide range of applications including classification, sequence modeling, and recommendation systems in real-world financial scenarios. The Orbis dataset [31] consists of quarterly and annual corporate revenue data across a variety of industries. It includes complementary features such as marketing expenditure, product launches, and industry trends, supporting use cases in benchmarking, revenue forecasting, and return on investment (ROI) analysis. The inclusion of macroeconomic indicators like interest rates and GDP enhances its utility for correlation and impact studies. Notably, the data is normalized across fiscal calendars, promoting consistency in comparative analytics. The CSMAR dataset [27] contains monthly-level financial records detailing cash inflows and outflows from operational, investing, and financing activities of firms varying in size and geography. It is purpose-built for studies on liquidity management, bankruptcy risk prediction, and working capital optimization. Annotated with financial events—such as credit line changes or tax deadlines—the dataset enables deeper temporal modeling and causality analysis in corporate finance research.

B. EXPERIMENTAL DETAILS

All experiments are conducted using PyTorch on NVIDIA A100 GPUs. We adopt the Adam optimizer with a learning rate of 1e-4 and apply a cosine annealing schedule for dynamic adjustment. Batch size is fixed at 64 for all datasets to ensure consistency. For training stability, we employ gradient clipping with a threshold of 1.0 and use mixed-precision training via AMP. Dropout with a rate of 0.1 is applied in all feedforward and attention layers to mitigate overfitting. Models are trained for 100 epochs with early stopping based on validation loss to prevent overfitting. We employ a standard train/validation/test split of 70%/15%/15% for all datasets. For temporal datasets such as Compustat and CSMAR, we strictly follow chronological order to avoid data leakage. Input sequences of length 12 (corresponding to 12 months or quarters, depending on the dataset) are used to predict the next step. For sequence modeling, we adopt a Transformer-based architecture with four encoder layers and a model dimension of 256. Positional encodings are learned rather than fixed to accommodate variable-length sequences. Layer normalization is applied after each sublayer, and residual connections are maintained throughout the network. Multi-head attention uses 8 heads with dimensionality 32 each. For classification tasks within the EDGAR dataset, we convert the task into a multi-class setting where each class corresponds to a transaction category. Cross-entropy loss is used as the objective function. In contrast, for regression tasks such as revenue or budget forecasting, we employ Mean Squared Error (MSE) loss. Mean Absolute Percentage Error (MAPE) and Root Mean Square Error (RMSE) are reported for comprehensive evaluation. All hyperparameters

TABLE 2. Evaluation of our method in contrast with current best-performing approaches on the Compustat and EDGAR datasets.

Model	Compustat Dataset				EDGAR Dataset			
	MAE ↓	RMSE ↓	MAPE ↓	R ² ↑	MAE ↓	RMSE ↓	MAPE ↓	R ² ↑
LSTM [31]	12.45±0.21	18.39±0.30	9.62±0.18	0.843±0.02	10.17±0.26	15.04±0.28	11.75±0.22	0.802±0.02
GRU [32]	11.98±0.25	17.86±0.33	9.33±0.21	0.851±0.01	10.63±0.24	14.68±0.30	10.88±0.24	0.815±0.01
Transformer [33]	11.50±0.20	16.77±0.27	8.97±0.19	0.869±0.02	9.94±0.22	14.21±0.26	10.24±0.20	0.827±0.02
DeepAR [34]	12.02±0.22	17.15±0.29	9.10±0.22	0.858±0.01	10.45±0.25	14.52±0.27	10.76±0.21	0.810±0.02
N-BEATS [35]	11.20±0.18	16.48±0.28	8.72±0.17	0.872±0.01	9.71±0.23	13.90±0.25	9.88±0.19	0.835±0.01
Informer [36]	10.93±0.19	16.25±0.26	8.50±0.18	0.878±0.02	9.60±0.21	13.84±0.24	9.72±0.18	0.840±0.01
Ours	9.84±0.16	14.39±0.22	7.29±0.15	0.902±0.01	8.45±0.18	12.77±0.20	8.12±0.16	0.869±0.01

TABLE 3. Comparative study of our approach and top-performing methods on the Orbis and CSMAR datasets.

Model	Orbis Dataset				CSMAR Dataset			
	MAE ↓	RMSE ↓	MAPE ↓	R ² ↑	MAE ↓	RMSE ↓	MAPE ↓	R ² ↑
LSTM [31]	13.21±0.26	18.44±0.34	10.02±0.22	0.836±0.02	11.15±0.24	16.82±0.31	11.95±0.23	0.808±0.02
GRU [32]	12.78±0.24	18.07±0.33	9.81±0.21	0.842±0.01	11.43±0.22	16.31±0.30	10.87±0.21	0.817±0.02
Transformer [33]	11.89±0.21	17.36±0.28	9.15±0.20	0.861±0.02	10.63±0.23	15.45±0.26	10.11±0.19	0.828±0.02
DeepAR [34]	12.34±0.22	17.79±0.30	9.30±0.19	0.853±0.01	11.02±0.25	15.82±0.27	10.44±0.20	0.814±0.01
N-BEATS [35]	11.56±0.19	16.95±0.27	8.89±0.18	0.867±0.01	10.28±0.22	15.11±0.24	9.90±0.18	0.837±0.01
Informer [36]	11.04±0.20	16.77±0.26	8.67±0.19	0.871±0.02	9.88±0.21	14.86±0.23	9.53±0.17	0.845±0.01
Ours	9.96±0.17	14.88±0.23	7.48±0.16	0.896±0.01	8.64±0.18	13.47±0.20	8.21±0.15	0.872±0.01

are selected based on a combination of grid search and prior work from top-tier conferences such as NeurIPS and ICLR. For baseline comparisons, we include classical statistical methods (ARIMA, VAR), traditional machine learning approaches (Random Forest, XGBoost), and recent deep learning models (LSTM, GRU, Transformer). Each baseline is carefully tuned to achieve its best performance under the same data split and evaluation metrics. Data preprocessing includes normalization via z-score transformation, missing value imputation using linear interpolation, and categorical embedding for all non-numeric metadata. We also apply feature selection using mutual information for datasets with high-dimensional inputs like Orbis. To evaluate generalization, we perform 5-fold cross-validation and report the average performance. All experiments are repeated with three different random seeds to ensure robustness, and the variance is reported along with the mean. The entire pipeline is implemented using reproducible code with fixed seeds, deterministic CUDA settings, and version-controlled dependencies.

C. COMPARISON WITH SOTA METHODS

We conduct a comprehensive comparison between our proposed method and several state-of-the-art (SOTA) baselines, across all four datasets. The results are presented in Table 2 and Table 3. We evaluate using four standard metrics including Mean Absolute Error (MAE), Root Mean Square Error (RMSE), Mean Absolute Percentage Error (MAPE),

and the coefficient of determination (R^2). Across all datasets and metrics, our method consistently outperforms existing models, showcasing its robustness and superior forecasting accuracy. on the Compustat dataset, our model achieves the lowest MAE of 9.84, outperforming Informer and N-BEATS by a significant margin. In parallel, our approach attains a 7.29 MAPE, demonstrating its effectiveness in reducing relative prediction error, especially compared to models like Transformer and GRU. Similarly, in the EDGAR dataset, our model achieves a high R^2 of 0.869, indicating strong fit and reliability in capturing temporal spending behavior. This suggests our method generalizes well to both structured fiscal planning data and fine-grained user transaction records. Our attention-based sequence aggregation and regularized temporal encodings offer advantages particularly in capturing hierarchical periodicity and suppressing outlier influence—characteristics common in both governmental and consumer financial data.

In Figure 5, the superiority of our model extends to the Orbis and CSMAR datasets. On Orbis, our model achieves an MAE of 9.96 and a R^2 of 0.896, significantly surpassing Transformer and Informer. Notably, while models like N-BEATS and DeepAR perform decently, they struggle to capture irregular seasonality and external drivers such as macroeconomic volatility. Our method, by incorporating adaptive feature weighting and recurrent residuals, effectively models both short-term surges and long-term trends. In the CSMAR dataset, where periodicity and irregular




FIGURE 5. Evaluation of our method in contrast with current best-performing approaches on the Compustat and EDGAR datasets.

Comparative Study: Performance Metrics on Orbis and CSMAR Datasets




FIGURE 6. Comparative study of our approach and top-performing methods on the Orbis and CSMAR datasets.

expenditures coexist, our method achieves the lowest MAPE of 8.21 and highest R^2 of 0.872. This result highlights the model’s capacity to balance inflow-outflow variations while minimizing forecast drift. Moreover, our improvement in RMSE across all datasets—particularly a 13.47 score in CSMAR—indicates reduced sensitivity to extreme errors, a common pitfall in financial domains. The result affirms our method’s resilience to variance and its ability to retain meaningful temporal features across different business cycles. These improvements are not trivial fluctuations but statistically robust gains consistently across folds and seeds.

In Figure 6, to further interpret the performance gain, we analyze the strengths of our model in light of its design choices and experimental results. Compared to Informer, which benefits from long sequence handling but suffers from static attention spread, our model applies time-decaying dynamic attention that prioritizes recent events while preserving historical context via gated residuals. DeepAR and LSTM often fail to generalize in regimes with abrupt changes or heterogeneous patterns due to their lack of adaptive reweighting, whereas our model benefits from learnable attention fusion and temporal bias correction modules. The advantage is especially evident in datasets like EDGAR, where transaction sequences are noisy and diverse. The superiority on Orbis can be attributed to our model’s hierarchical temporal block that captures quarterly and yearly variations simultaneously. Compared to Transformer, which uses fixed positional encodings, our model employs learnable positional

bias that adapts to seasonal shifts and irregular gaps, giving us a predictive edge. Moreover, callback to our method’s contribution—such as adaptive multiscale embedding and noise suppression—explains the robust performance across dynamic financial environments. These advantages culminate in consistent SOTA performance across all four benchmarks, validating our architectural innovations and demonstrating practical utility in real-world fiscal and financial forecasting.

D. ABLATION STUDY

To validate the effectiveness of each core component of our model, we conduct ablation experiments across all four datasets. The results are presented in Table 4 and Table 5, where we analyze the impact of removing individual modules including (Dynamic Graph Encoding) adaptive temporal attention, (Temporal-Aware Propagation) multi-scale feature encoder, and (Risk-Driven Policy Learning) regularized residual integration. From the results, it is evident that each component contributes significantly to the overall performance. The full model consistently achieves the best results across all metrics and datasets. For example, on the Compustat dataset, removing component Dynamic Graph Encoding leads to a notable degradation in R^2 (0.902 to 0.887) and increased MAE and RMSE, highlighting the critical role of our adaptive attention mechanism in learning dynamic temporal dependencies. Similarly, removal of Temporal-Aware Propagation or Risk-Driven Policy

TABLE 4. Evaluating module contributions in our framework using Compustat and EDGAR data.

Model	Compustat Dataset				EDGAR Dataset			
	MAE ↓	RMSE ↓	MAPE ↓	R ² ↑	MAE ↓	RMSE ↓	MAPE ↓	R ² ↑
w/o. Dynamic Graph Encoding	10.74±0.19	15.36±0.26	7.91±0.17	0.887±0.01	8.89±0.19	13.95±0.22	8.73±0.16	0.860±0.01
w/o. Temporal-Aware Propagation	10.18±0.18	14.97±0.24	7.55±0.15	0.891±0.01	8.61±0.17	13.69±0.21	8.28±0.15	0.864±0.01
w/o. Risk-Driven Policy Learning	10.43±0.20	15.11±0.25	7.78±0.16	0.888±0.01	8.74±0.18	13.88±0.22	8.45±0.15	0.861±0.01
Ours	9.84±0.16	14.39±0.22	7.29±0.15	0.902±0.01	8.45±0.18	12.77±0.20	8.12±0.16	0.869±0.01





FIGURE 7. Evaluating module contributions in our framework using Compustat and EDGAR data.

Learning leads to increased errors across both Compustat and EDGAR datasets, confirming that multi-scale encoding

and residual integration are complementary in capturing hierarchical and noisy patterns.

TABLE 5. Component evaluation of our framework using Orbis and CSMAR data.

Model	Orbis Dataset				CSMAR Dataset			
	MAE ↓	RMSE ↓	MAPE ↓	R ² ↑	MAE ↓	RMSE ↓	MAPE ↓	R ² ↑
w/o. Dynamic Graph Encoding	10.67±0.20	15.57±0.25	7.93±0.18	0.882±0.01	8.88±0.19	13.96±0.22	8.39±0.17	0.860±0.01
w/o. Temporal-Aware Propagation	10.21±0.19	15.01±0.24	7.60±0.17	0.888±0.01	8.74±0.18	13.74±0.21	8.19±0.15	0.865±0.01
w/o. Risk-Driven Policy Learning	10.43±0.18	15.25±0.23	7.77±0.16	0.884±0.01	8.69±0.17	13.82±0.22	8.33±0.16	0.863±0.01
Ours	9.96±0.17	14.88±0.23	7.48±0.16	0.896±0.01	8.64±0.18	13.47±0.20	8.21±0.15	0.872±0.01

**FIGURE 8.** Component evaluation of our framework using Orbis and CSMAR data.

In Figure 7, the performance drop observed in the Orbis and CSMAR datasets upon removing components Dynamic Graph Encoding, Temporal-Aware Propagation, or Risk-Driven Policy Learning further emphasizes the generalizability of our design. Without component Temporal-Aware Propagation, MAE increases from 9.96 to 10.21 and MAPE rises from 7.48 to 7.60 on Orbis, indicating the

TABLE 6. Robustness evaluation under high-noise scenarios (Gaussian noise $\sigma = 0.15$).

Model	MAE (↓)	RMSE (↓)	R ² (↑)
FinGraphNet	0.832	1.205	0.781
Transformer	0.964	1.372	0.692
GRU	1.016	1.461	0.644
LSTM	1.094	1.519	0.612
DeepAR	1.128	1.577	0.589

importance of multi-scale modeling in handling varying temporal granularity common in quarterly financial reports. Removing component Risk-Driven Policy Learning results in reduced R² from 0.896 to 0.884 and increased RMSE from 14.88 to 15.25, implying that regularized residual connections not only stabilize training but also capture long-range dependencies that simpler skip connections overlook. On the CSMAR dataset, each ablation results in consistent drops across all metrics. The model without adaptive attention (w/o. Dynamic Graph Encoding) shows the most significant error inflation, aligning with the fact that cash flow data often contains irregular bursts and lulls, which require adaptive temporal sensitivity to model accurately. This affirms that component Dynamic Graph Encoding is particularly vital in scenarios with non-stationary input sequences and volatile intervals.

In Figure 8, callback to the design motivations—component Dynamic Graph Encoding enables temporal relevance weighting based on learned context, component Temporal-Aware Propagation allows our model to disentangle overlapping temporal patterns through parallel convolutional encoders, and component Risk-Driven Policy Learning injects structured memory that preserves high-level trends without overfitting to short-term noise. These design choices result in a model that is not only accurate but also robust and interpretable across various financial scenarios. The performance gap between the full model and the ablated versions confirms that the strength of our method lies not in any single module but in the synergistic integration of all three. Together, these components offer complementary benefits including attention provides dynamic focus, multi-scale encoding improves representation diversity, and residual integration enhances temporal continuity. Therefore, the full architecture is essential for achieving state-of-the-art results across diverse and complex financial forecasting tasks.

As shown in Table 6, we evaluate the robustness of FinGraphNet and four baseline models under high-noise conditions, where Gaussian noise with a standard deviation of 0.15 is added to the input transaction features. The results demonstrate that FinGraphNet achieves the lowest MAE (0.832) and RMSE (1.205), while maintaining the highest R^2 score (0.781). In contrast, baseline models such as LSTM and GRU exhibit more significant performance degradation, with RMSE values exceeding 1.45 and R^2 dropping below 0.65. These findings suggest that FinGraphNet is more resilient to noise perturbations, likely due to its structure-aware graph encoding and constraint-regularized learning, which collectively enhance its stability in volatile financial environments.

TABLE 7. Sensitivity of model performance to the constraint penalty weight λ .

Penalty Weight (λ)	Avg. Reward	Constraint Violations	ROE (%)
0.1	143.2	24.5	12.1
0.5	150.7	18.3	13.6
1.0	156.4	11.7	14.3
2.0	152.8	9.5	13.9
3.0	141.3	5.2	11.2
5.0	132.6	3.6	10.3

In practice, enforcing hard constraints in a reinforcement learning setting can be problematic due to the need for feasible action sampling and policy gradient stability. Therefore, we adopt a soft penalty mechanism that incorporates constraint violations into the loss function using a squared hinge form. This approach enables the policy to learn near-feasible actions that gradually align with compliance boundaries without completely eliminating exploratory behavior. To validate this design choice, we conduct a sensitivity analysis on the penalty weight parameter λ in the constraint loss component. As shown in Table 7, we test values ranging from 0.1 to 5.0 and observe their impact on both reward and violation frequency. The results indicate that setting λ within [1.0, 2.0] achieves a favorable trade-off—producing compliant fiscal actions while maintaining strong performance. In contrast, values above 3.0 significantly reduce violations but at the cost of reward stagnation, illustrating the diminishing return of over-penalization.

V. CONCLUSION AND FUTURE WORK

In this study, we address the persistent challenges of adaptability and reliance on manual feature design in traditional accounting information classification systems. To overcome these issues, we propose a deep feature extraction framework based on convolutional autoencoders. This model is designed to learn hierarchical representations directly from raw accounting data, eliminating the need for handcrafted features. Our framework integrates joint objectives for data reconstruction and classification, enhanced by adversarial training to promote robustness against noisy and imbalanced entries. Experimental results on real-world accounting logs show that our method significantly improves classification

accuracy by 12 percentage points and anomaly detection F1-score by 15 percentage points compared to conventional shallow models. These findings highlight the efficacy of deep, automated feature learning in capturing complex patterns within accounting datasets, aligning with the growing trend of scalable, intelligent financial systems.

Despite the promising results, our method has several limitations. First, the convolutional-autoencoder architecture, while effective, may still struggle with extremely sparse or irregular financial data that lacks sufficient structure for convolutional learning. Second, the adversarial training introduces additional computational complexity and sensitivity to hyperparameter settings, which may hinder real-time deployment in high-frequency accounting environments. Third, while our graph-based temporal encoding is advantageous for modeling long-range dependencies, it involves pairwise interactions with quadratic time complexity ($O(n^2)$) in graph operations. This can become a computational bottleneck in large-scale enterprise systems. Fourth, the deep architecture may exhibit overfitting tendencies when applied to small-scale or low-diversity accounting datasets, limiting its generalization capability. In future work, we plan to explore hybrid models that combine symbolic reasoning with deep learning to better handle sparse data scenarios. We will also investigate efficient approximations for graph message passing, such as sampling-based or sparse attention mechanisms, to reduce graph operation overhead while preserving modeling capacity. Furthermore, we aim to incorporate regularization techniques and meta-learning strategies to mitigate overfitting risks, especially in low-resource settings. Another promising direction is to leverage domain-specific pretraining on large-scale financial corpora to improve the robustness and generalization of our framework across diverse accounting environments.

CONFLICT OF INTEREST STATEMENT

The author declares that the research was conducted in the absence of any commercial or financial relationships that could be construed as a potential conflict of interest.

ACKNOWLEDGMENT

This is a short text to acknowledge the contributions of specific colleagues, institutions, or agencies that aided the efforts of the author.

REFERENCES

- [1] R. M. Bushman and A. J. Smith, "Financial accounting information and corporate governance," *J. Accounting Econ.*, vol. 32, nos. 1–3, pp. 237–333, Dec. 2001.
- [2] B. Ballou, D. L. Heitger, and D. Stoel, "Data-driven decision-making and its impact on accounting undergraduate curriculum," *J. Accounting Educ.*, vol. 44, pp. 14–24, Sep. 2018.
- [3] U. Upadhyay, A. Kumar, G. Sharma, S. Sharma, V. Arya, P. K. Panigrahi, and B. B. Gupta, "A systematic data-driven approach for targeted marketing in enterprise information system," *Enterprise Inf. Syst.*, vol. 18, no. 8, Aug. 2024, Art. no. 2356770.
- [4] Z. Liu, "Accounting-oriented research on note recognition model based on information extraction algorithm," *WSEAS Trans. Bus. Econ.*, vol. 21, pp. 2640–2652, Dec. 2024.

- [5] S. Feng and R. Zhong, "Optimization and analysis of intelligent accounting information system based on deep learning model," *Comput. Intell. Neurosci.*, vol. 2022, Jul. 2022, Art. no. 1284289.
- [6] O. M. Lehner, C. Knoll, S. Leitner-Hanetseder, and C. Eisl, "The dynamics of artificial intelligence in accounting organisations," in *The Routledge Handbook of Accounting Information Systems*. Evanston, IL, USA: Routledge, 2022, pp. 121–139.
- [7] T. Sun, "Applying deep learning to audit procedures: An illustrative framework," *Accounting Horizons*, vol. 33, no. 3, pp. 89–109, Sep. 2019.
- [8] T. Sun, "Accounting information systems outputs: XBRL, AI and in-memory technologies," in *The Routledge Companion To Accounting Information Systems*. Evanston, IL, USA: Routledge, 2017, pp. 108–119.
- [9] J. M. Corchado, L. Borrajo, M. A. Pellicer, and J. C. Yáñez, "Neuro-symbolic system for business internal control," in *Proc. Ind. Conf. data mining*, 2004, pp. 1–10.
- [10] M. Hall, A. Ramsay, and J. Raven, "Changing the learning environment to promote deep learning approaches in first-year accounting students," *Accounting Educ.*, vol. 13, no. 4, pp. 489–505, Dec. 2004.
- [11] T. C. Redman, *Data Driven: Profiting From Your Most Important Business Asset*. Brighton, MA, USA: Harvard Bus. Press, 2008.
- [12] X. Zhang, "Application of data mining and machine learning in management accounting information system," *J. Appl. Sci. Eng.*, vol. 24, no. 5, pp. 813–820, 2021.
- [13] S. Cho, M. A. Vasarhelyi, T. Sun, and C. Zhang, "Learning from machine learning in accounting and assurance," *J. Emerg. Technol. Accounting*, vol. 17, no. 1, pp. 1–10, Mar. 2020.
- [14] L. Judijanto et al., "Data-driven marketing communications: The crucial role of accounting information systems," *Int. J. Financial econ.*, vol. 1, no. 2, pp. 511–520, 2024. [Online]. Available: <https://onlinelibrary.wiley.com/doi/abs/10.1002/cbdv.202402479>
- [15] S. Askary, N. Abu-Ghazaleh, and Y. Tahat, "Artificial intelligence and reliability of accounting information," in *Proc. Conf. E-Bus*. Cham, Switzerland: Springer, 2018, pp. 315–324.
- [16] P. O'Regan, *Financial Information Analysis: The Role of Accounting Information in Modern Society*. Evanston, IL, USA: Routledge, 2015.
- [17] S. Cohen, F. Manes Rossi, X. Mamakou, and I. Brusca, "Financial accounting information presented with infographics: Does it improve financial reporting understandability?" *J. Public Budgeting, Accounting Financial Manage.*, vol. 34, no. 6, pp. 263–295, Dec. 2022.
- [18] A.-D. Socea, "Managerial decision-making and financial accounting information," *Proc. Social Behav. Sci.*, vol. 58, pp. 47–55, Oct. 2012.
- [19] V. Chakraborty, V. Chiu, and M. Vasarhelyi, "Automatic classification of accounting literature," *Int. J. Accounting Inf. Syst.*, vol. 15, no. 2, pp. 122–148, Jun. 2014.
- [20] M. R. Garnsey, "Automatic classification of financial accounting concepts," *J. Emerg. Technol. Accounting*, vol. 3, no. 1, pp. 21–39, Jan. 2006.
- [21] A. Effiong, "Computerized accounting systems: Measuring structural characteristics," *Res. J. Finance Accounting*, vol. 11, no. 16, pp. 38–54, 2020.
- [22] S. Yang, "Strengthening accounting information systems with advanced big data mining algorithms: Innovative exploration of data cleaning and conversion automation," *Informatica*, vol. 49, no. 11, pp. 1–12, Jan. 2025.
- [23] C. A. Zhang, S. Cho, and M. Vasarhelyi, "Explainable artificial intelligence (XAI) in auditing," *Int. J. Accounting Inf. Syst.*, vol. 46, Sep. 2022, Art. no. 100572.
- [24] W. Xiuguo and D. Shengyong, "An analysis on financial statement fraud detection for Chinese listed companies using deep learning," *IEEE Access*, vol. 10, pp. 22516–22532, 2022.
- [25] P. Craja, A. Kim, and S. Lessmann, "Deep learning for detecting financial statement fraud," *Decis. Support Syst.*, vol. 139, Dec. 2020, Art. no. 113421.
- [26] O. V. Adamyk, "Difference between concepts 'automated', 'computer' and 'information' accounting systems: Transformation of elements of the method," *Ekonomichnyy analiz*, vol. 26, no. 1, pp. 163–169, 2016.
- [27] J. Li, C. Xu, B. Feng, and H. Zhao, "Credit risk prediction model for listed companies based on CNN-LSTM and attention mechanism," *Electronics*, vol. 12, no. 7, p. 1643, Mar. 2023.
- [28] J. J. Wild, *Financial Accounting: Information for Decisions*. New York, NY, USA: McGraw-Hill, 2017.
- [29] B. Ayash, Z. Bednarek, and A. Bordeman, "Compustat financial statement data articulation: A simplified approach," *J. Financial Reporting*, vol. 2025, pp. 1–18, Jun. 2025.
- [30] M. Choulga, G. Janssens-Maenhout, I. Super, E. Solazzo, A. Agusti-Panareda, G. Balsamo, N. Bousserez, M. Crippa, H. Denier van der Gon, R. Engelen, D. Guizzardi, J. Kuenen, J. McNorton, G. Oreggioni, and A. Visschedijk, "Global anthropogenic CO₂ emissions and uncertainties as a prior for Earth system modelling and data assimilation," *Earth Syst. Sci. Data*, vol. 13, no. 11, pp. 5311–5335, Nov. 2021.
- [31] J. Hromas, "Misalignment between profits and economic activity: Evidence from two distinct datasets," Tech. Rep., 2025.
- [32] S. M. Al-Selwi, M. F. Hassan, S. J. Abdulkadir, A. Muneer, E. H. Sumiea, A. Alqushaibi, and M. G. Ragab, "RNN-LSTM: From applications to modeling techniques and beyond—Systematic review," *J. King Saud Univ.-Comput. Inf. Sci.*, vol. 36, no. 5, Jun. 2024, Art. no. 102068.
- [33] S. Nosouhian, F. Nosouhian, and A. K. Khoshouei, "A review of recurrent neural network architecture for sequence learning: Comparison between LSTM and GRU," Tech. Rep., 2021. [Online]. Available: https://www.preprints.org/frontend/manuscript/3fa37c0d5d54bea69f5b855f25306b5e/download_pub
- [34] C. Generative Pre-Trained Transformer and A. Zhavoronkov, "Rapamycin in the context of Pascal's wager: Generative pre-trained transformer perspective," *Oncoscience*, vol. 9, pp. 82–84, Dec. 2022.
- [35] N. Schaduagrat, N. Anuwongcharoen, P. Charoenkwan, and W. Shoom-buatong, "DeepAR: A novel deep learning-based hybrid framework for the interpretable prediction of androgen receptor antagonists," *J. Cheminformatics*, vol. 15, no. 1, p. 50, May 2023.
- [36] B. N. Oreshkin, G. Dudek, P. Pelka, and E. Turkina, "N-BEATS neural network for mid-term electricity load forecasting," *Appl. Energy*, vol. 293, Jul. 2021, Art. no. 116918.
- [37] Q. Zhu, J. Han, K. Chai, and C. Zhao, "Time series analysis based on informer algorithms: A survey," *Symmetry*, vol. 15, no. 4, p. 951, Apr. 2023.
- [38] Y. Zhang, C. Chen, S. Zhu, C. Shu, D. Wang, J. Song, Y. Song, W. Zhen, Z. Feng, G. Wu, J. Xu, and W. Xu, "Isolation of 2019-nCoV from a stool specimen of a laboratory-confirmed case of the coronavirus disease 2019 (COVID-19)," *China CDC Weekly*, vol. 2, no. 8, pp. 123–124, Feb. 2020. [Online]. Available: <https://pmc.ncbi.nlm.nih.gov/articles/PMC8392928/>

FENGRI LIU received the bachelor's degree in accounting from Hunan University, Changsha, China, in 2022. From 2018 to 2024, he was an Auditor. Since 2023, he has been a Certified Public Accountant with Liaoning Dongding Certified Public Accountants, Benxi, China. He has participated in and led multiple financial auditing projects involving state-owned enterprises and listed companies. He has authored several articles in the fields of financial supervision, audit informatization, and enterprise risk management. His research interests include intelligent auditing systems, data-driven financial analysis, and deep learning applications in accounting practices. He is a member of the Chinese Institute of Certified Public Accountants. He actively contributes to regional forums on smart auditing transformation and has served as a peer reviewer for journals in accounting information systems.

...