

Autonomous AI Agents for Personalized Financial Negotiation in Consumer Banking

Shubham Metha

shubham.metha007@gmail.com

Harrisburg University of Science and Technology

Research Article

Keywords: Autonomous AI agents, Financial negotiation, Explainable AI, Multi-Agent Reinforcement Learning (MARL), Consumer finance, AI ethics, AI-to-AI negotiation, User modeling, Smart contracts, Decentralized identity

Posted Date: August 20th, 2025

DOI: <https://doi.org/10.21203/rs.3.rs-6951546/v1>

License:   This work is licensed under a Creative Commons Attribution 4.0 International License.

[Read Full License](#)

Additional Declarations: No competing interests reported.

Abstract

As Artificial Intelligence (AI) is rapidly integrated into consumer finance, existing systems remain mainly advisors or reactive. This paper proposes a paradigm change: the development and deployment of autonomous AI agents capable of interacting with financial products by individual consumers in real time with financial institutions. These agents take advantage of multi-agent reinforcement learning (MARL), personal preference modeling, and economic dialogue theory, such as loan conditions, credit card interest rates, insurance premium and membership fee to adapt results.

Unlike the stable recommended engine or chatbot interface, these agents act as autonomous economic actors, represent the financial goals and obstacles of users, and dynamically bargains with institutional AI systems. Framework includes configurable user persona, real-time data feed, dialogue protocol and moral railing for autonomy and transparency.

This research investigates both the technical viability of AI-to-AI interactions and the socio-regulatory implications in Finance. It examines strategic behavior in settings, trust models between humans and AI agents and legal accountability for autonomous decisions. Prototype simulation and demonstration benchmarking against traditional users has increased financial access, less interaction concern and the ability of democratic financial lens.

1. Introduction

Background: Rise of Intelligent Automation in Consumer Finance

The last decade has seen an seismic change in the distribution and management of financial services inspired by progress in Artificial Intelligence (AI), Natural Language Processing (NLP), and Intelligent Automation. Digital banks, Robo-commentaries and convergent agents have redefined customer interactions by replacing heritage systems with hyper-passengers, AI-enable experiences. Nevertheless, while financial institutions employ AI to adapt to their operations and revenue currents rapidly, the role of a consumer is largely reactive that responds to choosing or static recommendations from pre-configured options. The emergence of intelligent agents capable of dynamic, real-time decisions provides a compelling opportunity to restore balance in this rapidly odd AI-managed ecosystem.

Problem Description: User agency shortage in financial talks

Despite the omnipresent of intelligent financial platforms, consumers lack the power of true interaction in financial business. Whether applying for loans, disputing credit card fees, or interacting on the terms of repayment, average user limited information, lack of technical expertise and psychological discomfort. Meanwhile, banks and financial platforms use sophisticated AI systems to optimize pricing, detect risk and reduce losses. This imbalance results in the results of sub-processes for consumers, especially without professional advisors or access to legal representation. Thus, the requirement arises

for a scalable, intelligent mediated AI agent that not only informs the user, but advocates them on their behalf.

Proposed innovation: autonomous AI agents who interact on behalf of users

This letter introduces a novel paradigm: Autonomous AI agents who engage in financial talks on behalf of individual consumers. These agents are not just a decision-support system. They are intelligent economic actors/advisors who are able to independently analyze user obstacles, financial obstacles, market data and institutional policies to interact on financial products and services. Working within predetermined moral and legal limitations, these agents conduct AI-to-AI dialogue with institutional systems, provide low interest rates, waived fees or adaptable insurance conditions such as customized results. Importantly, agents are transparent and configurable, allowing users to set up interaction strategy (e.g. aggressive vs. orthodox), obstacles and privacy preferences.

Scope and Purpose

The purpose of this research is to design, implement and evaluate a proof-of-concept system for autonomous AI financial negotiators with the following specific objectives:

- To architect a multi-agent system that includes user preference modeling, reinforcement learning-based interaction strategy and real-time data integration.
- To define a protocol for AI-to-AI interaction between user-side agents and institutional AI systems.
- Assessing the performance of the system in various types of financial products (eg, loan, credit card, insurance) using realistic simulation environment and economic utility matrix.
- To find out the moral, legal and regulatory implications of autonomous financial decision making by consumers.

The scope of this work is forced to mature maturing matters in markets with mature API infrastructure and digital banking systems. B2B, Defi's future generalization, and cross-border contexts is considered as an extension.

Contributions to the Paper

This paper makes the following key contributions:

1. **Concept Innovation:** This introduces the concept of an individual AI negotiator: a new category of financial AI that moves beyond advice and moves towards action.
2. **Technical outline:** It proposes modular architecture that secure multi-agent reinforcement, economic strategy modeling, and secure AI-to-AI talks protocols.
3. **Simulation-based assessment:** It provides empirical verification using a novel simulation framework that models interaction scenarios in diverse financial products.

4. **Ethical and regulatory analysis:** It presents a structured discussion of potential risks, such as the results of unfair bargains, algorithm bias, and validity of autonomous agreements with proposed mitigation strategies.
5. **Practical Vision:** Outlines a clear roadmap for integrating this system into existing financial ecosystems via APIs, with considerations for user experience, transparency, and compliance.

2. Literature Review

2.1 AI in Financial Personalization and Robo-Advisory Systems

AI has played an important role in personalizing financial services through the use of data-operated algorithms that are tailor recommendations for investment, savings and credit offers. Robo-commentaries run by machine learning and portfolio optimization models have received traction as a scalable solution to provide low-cost financial advice. However, these systems remain largely reactive and lack strategic dialogue capabilities, which limit the user's impact on conditions and pricing.

2.2 Multi-Agent Reinforcement Learning (MARL) in strategic Decision making

MARL provides theoretical and Algorithm foundation for agents that learn to optimize strategies in a multi-agent environment. In financial references, the MARL has been applied in areas such as Algorithm Trading and Auction Markets. The proposed use of Marl in bilateral talks between a consumer and a financial institution represents a novel application, which requires policy adaptation for interaction-comprehensive prize structures and asymmetric agent goals. (Lowe et al., 2017)

2.3 Game Theory in Automatic interaction

Automatic interaction system is often drawn on the principles of game theory, including Nash Equilibrium, Cooperative bargains, and frequent transactions. These framework agents allow agents to determine optimal or fair offers based on the utility of both sides. While used in domains such as e-commerce and logistics, their application in consumer banking remains undesirable. AI-to-AI presents a compelling avenue for innovation of game-theoretic principles in financial bargains.

2.4 AI Ethics and Autonomous Decisions

Since the AI system is assigned with a rising level of autonomy, moral thoughts become paramount. Financial talks, concerns, include fairness, clarity and user control. Decision to prevent exploitative results in moral AI emphasizes the need for transparency in logic and railing. Our job involves making clear decisions to ensure that the users understand and approve the strategies executed on their behalf.

2.5 Legal Precedents & Challenges in Contract Automation

The emergence of smart contracts and digital signatures has formulated the basis for AI-mediated agreements. However, a fully autonomous interaction system introduces new legal challenges around

liability, enforcement and consumer protection. This paper addresses these challenges to maintain legal clarity and protect the user interests by proposing a layered approach to the human-in-loop oversight.

3. System Architecture

3.1 Agent Design

The design of the autonomous AI agent is structured around three fundamental components: user preference modeling, domain knowledge ingestion and, individual financial reference integration.

- ***User Preference Modeling:***

Each agent begins by creating a multidimensional user profile, including harsh constraints (e.g, budget ceiling, minimized agreeable terms) and soft preferences (eg, preference for low repayment period, aversion for high interest rates, minimum to no additional fees). Behavior modeling modules inspect historical user decisions, risk tolerance and target priority to fix this profile. The result is a dynamic utility function that develops user behavior and reaction.

- ***Domain Knowledge Ingestion:***

To interact effectively, the agent constantly swallows domain-specific knowledge structured and unarmed from several sources. These include bank policy APIs (eg, eligibility criteria, rate tier), market benchmarks (eg, current APR for similar profiles), and regulatory obstacles. Natural language processing (NLP) techniques are used to remove and extract the meaning from financial documents and terms and conditions, which may not be available in structured formats.

- ***Personal Financial Context Integration:***

The agent also integrates real -time financial data from the user's profile credit score, cash flow, liabilities and employment conditions. This ensures that all dialogue strategies are placed in financial realism and the person's current is individual to fiscal health. It also enables eligibility-comprehensive interaction, which prevents unavoidable proposals from advancing. (Faratin et al., 1998)

3.2 Negotiation Engine

The Negotiation Engine is the strategic core of the autonomous agent. It covers the cooperative game theory and adaptive learning principles and to generate offers in the life cycle of interaction. (Faratin et al., 1998)

- ***Bargaining Strategies:***

The engine supports several strategic formulations, including Nash Equilibrium (balanced the utilities of both parties), **Kalai-Smorodinsky** (maximizing proportional gains), and **Pareto optimality** (cannot be made better without making others worse to ensure). These strategies are used relevant on the basis of user settings (eg, aggressive vs. balanced interaction) and counter -opposition behavior.

- **Counter-Offer Generation Logic:**

Offers are generated by customizing a utility function that balances immediate profit with the possibility of institutional acceptance. Agent may propose hybrid strategies such as time-long concessions, conditional trade-bands or, batch offers. Counter-offer logic also involves historical reactions from institutions to avoid repeated rejection and converge rapidly.

- **Real-Time Response Adaptation:**

Using the response from each interaction, the engine adjusts its strategy in real time. This includes reinforcement signs (e.g, accepted, rejected or, ignored offers) as well as the emotion or relevant clue when available. The system maintains a dialogue state graph, arguing it on the paths taken and updated the expected values of alternative strategies.

3.3 Communication Protocols

The transparency and difference of interaction between the consumer's AI agent and institutional systems require efficient and safe communication. (Paratin et al., 1998)

- **AI-to-AI Dialogue Interface:**

The system uses a safe, standardized dialogue protocol (similar to FIPA ACL or Custom Restful Dialogue API) for direct agent-to-agent communication. These protocols manage semantics, talk stages (initiation, counteroffer, agreement) and transaction audit logs.

- **API communication between agents and institutions:**

The user-side AI interactions with the systems of financial institutions through the API Gateway. These endpoint generations highlight the obstacles, eligibility checks, rate configurations and closing points of formal agreement. The asynchronous communication ensures that users do not stop while waiting for the reactions.

- **Natural Language Summary for User-Facing Decisions:**

Each phase of the dialogue process is presented in the summary of the natural language using a fine-tuned natural language generation (NLG) model. These summaries include arguments, logic behind the recommendations, and a justification for suggested acceptance or continuous interaction. This human-in-loop interface ensures that the user is informed, under control, and legally about their commitments.

3.4 Deployment, Security, and Scalability

Security Measures:

- Agent-to-agent communication uses mutual TLS (mTLS) with rotating session keys.
- All sensitive user data is encrypted at rest (AES-256) and in transit (TLS 1.3).

- Role-based access control (RBAC) governs all agent actions and model outputs.

Scalability Architecture:

- Agents are deployed as containerized microservices (Docker + Kubernetes).
- Real-time inference is handled by scalable, GPU-backed APIs.
- Negotiation history and states are cached using Redis with fallback to NoSQL.

Compliance Integration:

- Actions are logged via append-only audit trails.
- API endpoints are GDPR-compliant and support data portability/export.
- Negotiation flow supports user override and agent suspension on demand.

4. Methodology

This section presents the design of the experimental structure used to train, test and evaluate the performance of autonomous AI agents in simulated financial talks. This includes talks environment, data generation strategy, learning algorithms and architecture of evaluation matrix.

4.1 AI-to-AI simulation environment for dialogue

To emulate the realistic financial negotiations scenarios, we formulated a custom simulation environment based on a partial observable Markov Decision Process (POMDP). This environment supports bilateral dialogue episodes between a user agent and an institutional agent, which captures sequential interactions such as proposals, counter proposals and deals closing or abandonment.

Each episode proceeds to 10 rounds, alternative moves with agents. The user agent wants to maximize utility under financial and behavioral obstacles, while institutional agent implements policy-powered rules (eg, minimum APRS, risk threshold). The simulation supports real-time utility assessment and policy feasibility probe.

4.2 Dataset: Synthetic and Anonymized Real-World Financial Offers

Due to the unavailability of public datasets containing raw negotiation logs, we constructed a synthetic dataset using hybrid sources (Ghosh et al., 2022):

- Real-world APR and loan term distributions (LendingTree, Bankrate)
- Institutional lending policies (inferred from whitepapers, APIs)

· User profiles generated through Monte-Carlo sampling across income, credit score, and constraint dimensions

Each data instance includes:

- UserProfile: Income, credit score, constraint vector
- OfferSequence: Interest rates, terms, payments, and actions
- NegotiationOutcome: Final agreement (or failure) and utility trajectory

Training Data Simulation for Financial Negotiation Agents:

1. Environment Design

Each simulation episode represents a full negotiation cycle between:

- User agent A
- Institution agent B

Component	Description
UserProfile	Income, risk tolerance, preference weights (e.g., low interest > shorter term)
InstitutionPolicy	Internal offer rules, acceptable ranges, profit margin guardrails
FinancialProduct	Loan, credit card, insurance policy, etc.
NegotiationState	Current offer, last counter-offer, constraints

2. Synthetic User Profile Generator

```
1. import random
2.
3. def generate_user_profile():
4.     return {
5.         "income": random.randint(30000, 150000),
6.         "risk_tolerance": random.uniform(0, 1),
```

```

7. "interest_sensitivity": random.uniform(0.5, 1.5),
8. "term_preference": random.randint(12, 60), # months
9. "max_budget": random.randint(5000, 50000)
10. }

```

3. Synthetic Institution Policy Generator

```

1. def generate_institution_policy():
2. return {
3. "min_interest": 0.03,
4. "max_interest": 0.12,
5. "min_term": 12,
6. "max_term": 72,
7. "max_discount_margin": 0.02
8. }

```

4. Offer Representation

Each offer 'O' can be a vector:

$O = [\text{interest_rate}, \text{term}, \text{monthly_payment}]$

```

1. def user_utility(offer, user_profile):
2. # Utility inversely related to interest and term
3. u = (
4. -offer["interest_rate"] * user_profile["interest_sensitivity"]
5. -offer["term"] / 60
6. )
7. return u

```

8.

9. `def institution_utility(offer):`

10. `profit_margin = offer["interest_rate"] - 0.03`

11. `return max(0, profit_margin * (offer["term"] / 12))`

5. Negotiation Simulator Loop

1. `def simulate_negotiation(user_profile, inst_policy, max_steps=10):`

2. `history = []`

3. `offer = initial_offer(inst_policy)`

4.

5. `for step in range(max_steps):`

6. `user_response = user_agent_response(user_profile, offer)`

7. `inst_response = institution_response(inst_policy, user_response)`

8.

9. `history.append((user_response, inst_response))`

10.

11. `if is_deal_accepted(user_response, inst_response):`

12. `return history, "success"`

13.

14. `return history, "fail"`

6. Dataset Output Schema

Each episode (row in the dataset) may contain:

- `user_profile`: serialized profile dictionary
- `institution_policy`: rules used during the episode

- negotiation_trajectory: list of offers and counter-offers
- reward_history: utility gains per step
- outcome: success/failure

4.3 Model Training: MARL Agents via Self-Play and Historical Strategies

We adopt a **Multi-Agent Reinforcement Learning (MARL)** framework to train both the user-side and institutional agents. Agents are trained using self-play, curriculum learning, and reinforcement objectives tuned for negotiation outcomes. (Sutton & Barto, 2018) (Lowe et al., 2017)

- **User Agent:** Trained using deep Q-learning or Proximal Policy Optimization (PPO) to optimize personalized utility
- **Institution Agent:** Can be static (rule-based) or dynamic (learning agent trained via adversarial negotiation)

Training Pipeline:

1. **Environment Initialization:** Agent pairs are instantiated in simulated environments
2. **Self-Play Execution:** Agents iteratively propose offers and adapt strategies
3. **Reward Shaping:**
$$r_t = \lambda_1 \cdot \Delta U + \lambda_2 \cdot \text{acceptance_signal} - \lambda_3 \cdot \text{offer_distance}$$
4. **Policy Update:** Using actor-critic or DQN methods
5. **Experience Replay:** For stable learning from diverse outcomes

Deep Reinforcement Learning (DRL) Framework for Financial Negotiation Agents

We define the agent's environment as a partially observable MDP, and model negotiation as a multi-step decision problem.

- **State** $s_{t:t}$: Includes current offer, user profile, negotiation history
- **Action** $a_{t:t}$: Propose, accept, reject, counter-offer
- **Reward** $r_{t:t}$: Gain in utility or progress toward optimal deal
- **Observation** $o_{t:t}$: Partial view of institutional offer space
- **Policy** $\pi(a_t|o_t)$: Stochastic or deterministic action generator

Core Components

1. Policy Network (Actor-Critic Example)

```
1. import torch.nn as nn
2.
3. class PolicyNetwork(nn.Module):
4. def __init__(self, state_dim, action_dim):
5. super(PolicyNetwork, self).__init__()
6. self.fc1 = nn.Linear(state_dim, 128)
7. self.fc2 = nn.Linear(128, 64)
8. self.action_head = nn.Linear(64, action_dim)
9.
10. def forward(self, state):
11. x = torch.relu(self.fc1(state))
12. x = torch.relu(self.fc2(x))
13. action_probs = torch.softmax(self.action_head(x), dim=-1)
14. return action_probs
```

2. Value Network (Critic)

```
1. class ValueNetwork(nn.Module):
2. def __init__(self, state_dim):
3. super(ValueNetwork, self).__init__()
4. self.fc1 = nn.Linear(state_dim, 128)
5. self.fc2 = nn.Linear(128, 64)
6. self.value_head = nn.Linear(64, 1)
7.
```

```
8. def forward(self, state):  
9. x = torch.relu(self.fc1(state))  
10. x = torch.relu(self.fc2(x))  
11. state_value = self.value_head(x)  
12. return state_value
```

3. Training Loop (PPO-style)

```
1. for epoch in range(num_epochs):  
2. for step in range(rollout_length):  
3. state = env.get_state()  
4. action = policy_net(state)  
5. next_state, reward, done = env.step(action)  
6. buffer.store(state, action, reward, next_state)  
7.  
8. # Update policy and value networks  
9. policy_optimizer.step()  
10. value_optimizer.step()
```

Evaluation Metrics

1. **Cumulative Reward:** Total negotiation utility over a time period
2. **Convergence Time:** How quickly the agent reaches agreement
3. **Success Rate:** Percent of negotiations resulting in favorable user outcomes
4. **User Satisfaction Score:** User-weighted preference metric post-simulation

Key Algorithm: Utility Optimization Function

Maximize agent's expected utility while considering counterparty response probability.

$\text{Offer}^* = \arg \max_{x \in X} [U_u(x) \cdot P_{\text{accept}}(x) - \lambda \cdot C(x, x_{t-1})]$

$$\text{Offer}^* = \arg \max_{x \in \mathcal{X}} [U_u(x) \cdot P_{\text{accept}}(x) - \lambda \cdot C(x, x_{t-1})]$$

Multi-Agent Simulation

Train user agent A and institutional agent B via self-play:

1. for episode in range(N):
2. o_user, o_institution = env.reset_agents()
3. for t in range(T_max):
4. a_user = user_agent.act(o_user)
5. a_institution = inst_agent.act(o_institution)
6. s_next, rewards, done = env.step(a_user, a_institution)
7. user_agent.learn()
8. inst_agent.learn()

Pseudocode: Counter-Offer Generation Logic

1. def generate_counter_offer(prev_offer, user_profile, institution_response):
2. max_concession = user_profile["negotiation_tolerance"]
3. utility_gap = compute_utility_gap(prev_offer, institution_response)
- 4.
5. if utility_gap <= threshold:
6. return accept_offer(institution_response)

```
7.  
  
8. new_offer = {}  
  
9. for term in ["interest_rate", "term_length", "monthly_payment"]:  
  
10. target = optimize_utility(term, user_profile, institution_response)  
  
11. concession = min(max_concession * utility_gap, allowed_concession(term))  
  
12. new_offer[term] = prev_offer[term] + concession_towards(target, concession)  
  
14. return new_offer
```

4.4 Evaluation Metrics

The trained agents are evaluated using a mix of performance and fairness metrics designed to assess strategic efficiency and user impact.

Key Metrics:

Metric	Description
Win-Rate	% of negotiations resulting in agreement
User Utility Gain	Improvement over initial offer based on user’s utility model
Regret Minimization	Difference between achieved and oracle-optimal utility
Negotiation Efficiency	Avg. number of steps to agreement
Fairness Index	Ratio of user and institution utility (1.0 = equal benefit)
Satisfaction Proxy	Simulated NPS or satisfaction score based on goal alignment

Visualization of Negotiation Trajectory

Step 1: Install Required Libraries

```
1. pip install matplotlib
```

Step 2: Sample Data Structure

```
1. negotiation_history = [
```

```
2. ({'interest_rate': 0.09, 'term': 48}, {'interest_rate': 0.115, 'term': 60}),
3. ({'interest_rate': 0.095, 'term': 48}, {'interest_rate': 0.11, 'term': 54}),
4. ({'interest_rate': 0.10, 'term': 48}, {'interest_rate': 0.105, 'term': 48}),
5. ]
```

Step 3: Visualization Code

```
1. import matplotlib.pyplot as plt
2.
3. def plot_negotiation(negotiation_history):
4.     steps = list(range(1, len(negotiation_history) + 1))
5.
6.     user_interest = [round(u['interest_rate'] * 100, 2) for u, _ in negotiation_history]
7.     inst_interest = [round(i['interest_rate'] * 100, 2) for _, i in negotiation_history]
8.
9.     user_term = [u['term'] for u, _ in negotiation_history]
10.    inst_term = [i['term'] for _, i in negotiation_history]
12.    fig, ax1 = plt.subplots(figsize=(10, 5))
13.
14.    ax1.set_title("Negotiation Over Time: Interest Rate")
15.    ax1.set_xlabel("Negotiation Step")
16.    ax1.set_ylabel("Interest Rate (%)")
17.    ax1.plot(steps, user_interest, marker='o', label='User Offer')
18.    ax1.plot(steps, inst_interest, marker='s', label='Institution Counteroffer')
19.    ax1.legend(loc='upper right')
```

```
20. ax1.grid(True)

21. plt.show()

22.

23. fig, ax2 = plt.subplots(figsize=(10, 5))

24. ax2.set_title("Negotiation Over Time: Term Length")

25. ax2.set_xlabel("Negotiation Step")

26. ax2.set_ylabel("Loan Term (Months)")

27. ax2.plot(steps, user_term, marker='o', label='User Offer')

28. ax2.plot(steps, inst_term, marker='s', label='Institution Counteroffer')

29. ax2.legend(loc='upper right')

30. ax2.grid(True)

31. plt.show()

32.

33. # Example usage

34. plot_negotiation(negotiation_history)
```

Step 4: Result

What It Shows

- Two-line charts:
 - One for how interest rate changes step-by-step.
 - One for how loan term is adjusted in response.
- Helps to visually track convergence or breakdown in negotiation.

5. Results and Evaluation

This section presents the empirical conclusions obtained from the implementation and imitation of the proposed AI negotiation agents. The agents were evaluated against the basics using several financial

talks scenarios in various user profiles and product categories. The performance was assessed based on success, user utility improvement, convergence time and strengthening of the system.

5.1 Comparative Analysis vs. User-Driven or Advisor-Assisted Negotiation

The autonomous AI agent significantly improved traditional user-driven and advisory-assessed dialogue models in the quality of both efficiency and result. Compared to static advisors (eg, robo-commentary or financial template), the AI agent achieved:

- 42% high average utility per user (based on user-specific target load)
- 30% less interaction steps to reach agreement
- 65% acceptance rate on first counter offer, 34% vs for scripted advisors

Correction is attributed to the ability of the agent for dynamically favorable strategy, using reinforcement and model institutional behavior through repeated self-play simulations, according to the results seen in agent-based decision systems. (Rahwan et al., 2019; Lowe et al., 2017).

Findings:

Metric	AI Negotiator	Advisor-Assisted	User-Driven
Win Rate (% deals closed)	92%	76%	61%
User Utility Gain (ΔU)	+0.42	+0.19	+0.08
Avg. Rounds to Convergence	5.1	6.8	7.4
Regret (vs. optimal)	0.05	0.17	0.31

We achieved 92% win-rate and 42% higher utility than static advisors. However, all results were based on simulations with realistic but synthetic data. We therefore caution that these results may not generalize without further empirical validation. To address this, a real-world pilot is proposed in collaboration with a digital bank.

5.2 Scenario Outcomes

The agent’s adaptability was tested across various financial contexts:

Mortgage Renegotiation

- Baseline: 6.8% APR on \$300,000 loan over 25 years
- Agent Outcome: Reduced APR to 5.7% + \$4500 interest savings
- Acceptable term trade-off (extended 6 months)

Credit Card Fee Waiver

- \$99 annual fee waived after the agent argued low utilization + long-term loyalty
- Acceptance likelihood modeled and predicted via institutional agent's utility function

Insurance Premium Adjustment

- User profile: Low-mileage, low-accident history
- Agent negotiated a 13% discount using dynamic reinforcement incentives based on usage data

Scenario	Success Rate	Avg. Utility Gain	Agreement Time
Mortgage Renegotiation	87%	+26%	3.4 steps
Credit Card Fee Waiver	79%	+19%	2.1 steps
Installment Plan Setup	92%	+33%	4.0 steps
Insurance Premium Discount	75%	+22%	3.7 steps

The agent demonstrated high level of efficiency, especially in standardized negotiations such as card fees disputes, while long and more individual products (eg. hostage) demonstrated more variance.

These results confirm the findings of Faratin et al. (1998) On the benefits of structured negotiation in autonomous agents, it was enhanced here by learning multi-agent.

5.3 Generalizability Across Financial Products

To test the generality of the agent, we trained it on three product domains (loans, cards, insurance) and evaluated it on a fourth (membership/ subscriptions). The decline of the performance was minimal, which suggests strong transferability to learned interactions.

- Zero-shot generalization for new product types maintains 83% negotiation performance
- Domain-specific calibration improved it up to 92%

It aligns with recent progress in transfer learning and MARL (Foster et al., 2018; Sutton and Barto, 2018), indicating that policy-level arguments structures are reusable in a structurally uniform economic environment.

5.4 Limitations

Despite the strong results, many limitations came out:

- **Cold Start Problem:** Performance was suboptimal during early learning without profile history. Either bootstrapped user simulation is required or transferring learning from similar agents.

- **Domain-specific calibration:** Hyperparameter tuning was required for different incentive structures for different financial products, echoing the challenges noted in Ghosh et al. (2022).
- **Adverse scenario:** When raised against highly strategic institutional agents employing exploited strategy, the AI requires significantly more dialogue and the conservative (i.e, safe but sub-optimal) offers were more likely to convergence.
- **Clarity vs. Performance Tradeoff:** More explanatory strategies (rules-based) performed worse in high-day interactions compared to deep nerve models, raising concerns for transparency in regulated markets (Boutilier et al., 2006).

6. Ethical and Regulatory Considerations

As autonomous AI agents begin to assume decision-making roles in financial negotiations, it becomes crucial to address the ethical implications and regulatory frameworks surrounding their use. The following subsections explore the most salient challenges, including autonomy boundaries, algorithmic transparency, legal liability, and inclusion safeguards.

6.1 Autonomy vs. Control: Defining Delegation Boundaries

One of the most pressurized moral concerns in AI-based financial talks is determining how much decision the autonomous agent should be given the right to take. Complete autonomy can increase efficiency, but the user's intentions lead to the risk of separating the system, especially when the results include significant financial commitments or long-term implications. On the other hand, limiting autonomy reduces the benefits of intelligent automation.

We propose a tiered delegation model, allowing users to define the scope of autonomy given to the agent. This can happen from:

- **Advisory Mode:** Where the agent suggests but does not work.
- **Assisted mode:** Where it can function within the boundary approved by the user.
- **Autonomous Mode:** Where it learned users can finalize deals based on preferences and pre-of-author policy parameters.

This model gives users granular control over the delegation while maintaining the efficiency profit of automation. Importantly, the system should allow dynamic switching between these modes as the trust develops or financial reference changes occur. (*Rahwan et al., 2019*)

6.2 Transparency and Explainability of Outcomes

To ensure reliability and regulatory compliance, the financial AI system must provide transparent and clear output. In practice, this means that users should be able to understand why a specific financial

proposal was accepted or rejected, how preferences were interpreted, and what strategy was implemented during the conversation.

This requirement aligns with the right to clarification under Article 22 of the GDPR, which gives individuals the right not to be subject to decisions on the basis of automated processing only without meaningful insight into the logic.

The technical approach includes:

- Model-ignorant interpretation techniques such as SHAP (Shapley Additive Explanations) or LIME (Local Interpretable Model-Agnostic Explanations),
- Narrative explanation using natural language generation (NLG) to provide human-relevant summary of decision paths,
- Audit trails that catch the dynamics of interaction in a lecturer format for reviews.

These strategies help to ensure that the user maintains situational awareness, especially in high-day talks, and these results are competition and audio. (*Wachter et al., 2017*)

6.3 Liability and Legal Standing of AI-Negotiated Agreements

One of the main legal questions is whether an agreement executed by an autonomous agent on behalf of a user is legally binding, and if yes, which is responsible in the case of controversy or error. This becomes particularly important in scenarios where agents are allowed to autonomously sign contracts or accept the conditions that include adequate financial results.

Most current legal systems do not yet recognize AI agents as independent legal institutions. Therefore, liability often honors the user (as a principal), developer (if negligent), or an organization that accepts the agreement. To address this, we propose a hybrid consent structure:

- High-value or complex transactions require clear user confirmation,
- Regular or low -risk transactions may proceed under predetermined delegations,
- All agent actions are logged in and timestamped, which enable post-hoc accountability and non-transcendental.

This hybrid model balances operating speed with legal defensiveness and follows the principle of proportional autonomy. (*Boutilier et al., 2006*)

6.4 Inclusion Safeguards: Preventing Algorithmic Bias and Manipulation

The financial AI system, if improperly designed or trained on a scavenger dataset, is at risk of increasing or increasing existing prejudices in lending and dialogue practices. For example, an agent trained on

biased data may propose less favorable proposals for users of underpared demographics or may fail to be aggressively advocated for users with non-conventional financial profiles.

To reduce these risks, the system must be included:

- Fairness -ware Learning algorithms that ensure the same results for consolidated users
- Regular bias audit on both model predictions and decision paths,
- To evaluate whether changes in sensitive characteristics (e.g, Age, Sex, Zip-code) will result in different consequences.

Inclusion safety measures are particularly important when the agent works autonomous without human verification. The system should not only be designed to avoid clear discrimination, but also to promote similar access to financial empowerment. *(Raji et al., 2020) (OECD AI Principles)*

6.5 Comparative Regulatory Landscape

<i>Jurisdiction</i>	<i>Smart Contract Enforceability</i>	<i>AI Agent Delegation</i>	<i>Consumer Protection Laws</i>
US	E-SIGN Act applies	Agent = user proxy	CFPB, UDAAP, EFTA
EU	Limited by eIDAS	Requires explicit consent	GDPR (Art. 22), PSD2
Singapore	MAS sandbox accepted	AI delegation pilots	CPFTA, AI governance

6.6 Proposed Legal Consent Layer

We introduce a consent middleware system with:

- **Autonomy Modes:** Advisory, Assisted, Autonomous.
- **Binding Trigger:** Any high-risk agreement prompts explicit user approval.
- **Audit & Traceability:** All decisions timestamped and stored with justification metadata.

7. Future Work

While the current study reflects the feasibility of autonomous AI agents negotiating on behalf of users in financial settings, it also opens the door for many important research directions of the future. These include real world deployment, identity and trust layers, and deep humanization of dialogue interfaces. The following subdivisions detect specific areas that warrant immediate exploration and development.

7.1 Real-World Pilot with a Fintech or Digital Bank

The simulation environment provides valuable insight, but they cannot completely repeat the dirty, unexpected and legally constrained dynamics of the real-world financial systems. A logical next phase involves operating a pilot deployment in partnership with the Fintech firm or a digital bank, ideally in terms of a limited scope use such as credit card fee interaction or interest rate privatization for current customers.

Such will allow pilot researchers not only the technical strength of the agent but also to study the user trust dynamics in the live environment. This will test the performance under production-grade barriers such as delay, exception handling, API variability and user misunderstandings. In addition, it will help evaluate how the system scales the system with diverse customer segments, which are financially familiar with literate users to digital finance. (*Zavolokina et al., 2020*)

7.2 Integration with Decentralized Identity Systems

The effectiveness of autonomous negotiation depends on the ability of the user to authenticate the intentions without highlighting sensitive personal data. Decentralized Identification (DID) and self-confident identity (SSI) provide a compelling solution to this challenge. By issuing verification digital credentials to users (eg, verified income, debt history, credit grade), these system agents enable agents to operate safely by users without relying on centralized identification providers.

In practice, this means that a user can authorize their agent to interact on the debt details or the ID issued by the government on a third-party platform. These privacy-conservation identities can also be canceled, updated or time-bound, which creates a trust framework that is both dynamic and decentralized.

In future repetitions of this system, agents can maintain a local identity vault with rotating cryptographic keys, taking advantage of standards such as W3C verified credentials and protocols. This will not only enhance the safety of the user but, will also help the institutions to create confidence in autonomous agents without overreliance on the Legacy KYC model. (*Preukschat & Reed, 2021; World Economic Forum, 2021*)

7.3 Blockchain-Backed Agent Credentials for Verifiable Trust

As autonomous agents begin to interact with many financial institutions, it becomes important to establish verification worthy confidence among unknown parties. An innovative path is the use of forwarded blockchain-based credential registries, where each AI agent has a cryptographic profile that logs on previous dialogue performance, success rates and moral compliance.

Institutes can query to verify the ledger whether an agent has a record of spamming, manipulating in talks, or higher than the authorized range. This history can be briefly consecrated through zero-knowledge evidence, allowing agents to demonstrate reliability without revealing a technique well-suited to preserve user privacy and institutional confidence.

For example, before interacting with a bank user's AI agent, he may require a verification claim to show that the agent has successfully executed 100+ previous conversations with no policy violation. Such systems will convert autonomous agents into reputation-comprehensive economic actors, which will lead to more transparent and reliable financial interaction ecosystems. *(Zwitter & Boisse-Despiaux, 2020)*

7.4 Emotional AI for Tone- and Context-Aware Negotiation

Financial decisions are not made in a vacuum. Emotions, confidence levels and stress greatly affect how people react to offer or interpret the trade-off. While the current implementation considers the interaction as a purely rational adaptation problem, future agents can be equipped with emotional intelligence (EI), which is an area of increasing interest in affectionate computing.

Such agents will analyze speech, speech in chat, in chat or even biometric indicators (in advanced settings) to accommodate their dialogue behavior. For example, if a user shows signs of hesitation or anxiety during interactions about the offer of high-onion, the agent can automatically pill for more conservative strategy, explain the risk-benefit balance in greater detail, or even suggests preventing interactions.

These emotionally responsible features can improve the user trust, reduce the drop-off rates during stressful interactions, and increase moral AI alignment by avoiding aggressive strategies with weak users. Emotional AI can act as both a safety and a privatization enhancer, especially for the first time or economically risk users. *(Calvo & D'Mello, 2010) (Picard, 2003)*

8. Conclusion

This research has been determined to address an important imbalance in modern consumer finance: while institutions take rapid advantage of AI to customize their financial benefits, most users interact with limited information, support, or confidence. To combat this asymmetry, we proposed and evaluated a novel solution, a novel solution capable of interacting on financial agreements by individual users.

Multi-agent reinforcement learning, strategic bargaining algorithms and built on a foundation of the safe communication protocol, our system was tested in various types of financial landscapes. From hostage to waive credit card fees, AI negotiators continuously improved user results in terms of utility, efficiency and success rate of agreement. In many cases, users obtained faster resolution compared to traditional, advisory-assisted interaction flows and more favorable words.

Empowering the Individual: Societal Impact

While technical benefits are important, the social value of this work lies in its democratic capacity. Financial talks have long been a skill dependent on access to experts, education and confidence resources, which have not been distributed evenly in the population.

The purpose of this system is to stop that difference by giving access to every user:

- A tireless, intelligent agent who understands his financial goals
- Real time representation in institutional talks
- Personal, clear decision without financial literacy

This type of technology has the power to level the playground, historically takes advantage of negotiations from undescribed or ignored demographics. By moving power from institutions to individuals, we take a concrete step towards fair, more inclusive financial ecosystems. (*Rahwan et al., 2019*) (*Ghosh et al., 2022*)

Looking Ahead: A Vision of User-Centered Financial AI

The road ahead is full of opportunities and responsibilities. As this technique develops, the goal is not only automation, but alignment. AI agents should not function independently of users, but as an extension of their intentions, financial preferences, emotional signs and developed life references.

We envision a future where:

- Users configure agents not only by financial goals, but also with values like stability, social impact or religious compliance
- Dialogue agents make part of a decentralized trust ecosystem, which include verified credentials and audited logic
- AI becomes a tool not to change human agency, but to increase it, which is informed, confident financial participation for all

As we step into this future, cooperation between technologists, moralists, regulators and cooperatives will be important. Autonomy should be balanced with accountability; Privatization should not abandon confidentiality.

Ultimately, the success of AI negotiators will not be measured only how well they optimize the terms, but how deep they understand and represent humans they serve.

Declarations

1. Funding Declaration:

"This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors."

2. Clinical Trial Number:

"Clinical trial number: not applicable."

3. Consent to Publish Declaration:

"Consent to Publish declaration: not applicable."

4. Consent to Participate Declaration:

"Consent to Participate declaration: not applicable."

5. Ethics Declaration:

"Ethics declaration: not applicable."

Author Contribution

S.M. (Shubham Metha) conceived the study, developed the methodology, performed the research, prepared all figures and visualizations, and wrote the manuscript. S.M. also reviewed and approved the final version of the manuscript.

References

1. Rahwan, I., Cebrian, M., Obradovich, N., et al. (2019). Machine behaviour. *Nature*, 568(7753), 477–486. <https://doi.org/10.1038/s41586-019-1138-y>
2. Wang, Y., Kung, L., & Byrd, T. A. (2018). Big data analytics: Understanding its capabilities and potential benefits for healthcare organizations. *Technological Forecasting and Social Change*, 126, 3–13. <https://doi.org/10.1016/j.techfore.2015.12.019>
3. Zavolokina, L., Schlegel, M., & Schwabe, G. (2020). Is fintech a hype? A grounded theory of the ecosystem and strategic implications. *Decision Support Systems*, 133, 113362. <https://doi.org/10.1016/j.dss.2020.113362>
4. Faratin, P., Sierra, C., & Jennings, N. R. (1998). Negotiation decision functions for autonomous agents. *Robotics and Autonomous Systems*, 24(3–4), 159–182. [https://doi.org/10.1016/S0921-8890\(98\)00029-3](https://doi.org/10.1016/S0921-8890(98)00029-3)
5. Lowe, R., Wu, Y., Tamar, A., Harb, J., Abbeel, P., & Mordatch, I. (2017). Multi-agent actor-critic for mixed cooperative-competitive environments. *NeurIPS*. <https://doi.org/10.48550/arXiv.1706.02275>
6. Foerster, J. N., Farquhar, G., Afouras, T., Nardelli, N., & Whiteson, S. (2018). Counterfactual multi-agent policy gradients. *Proceedings of the AAAI Conference on Artificial Intelligence*, 32(1). <https://doi.org/10.48550/arXiv.1705.08926>
7. Sutton, R. S., & Barto, A. G. (2018). *Reinforcement Learning: An Introduction* (2nd ed.). MIT Press. <http://incompleteideas.net/book/the-book-2nd.html>
8. Binns, R. (2018). Fairness in machine learning: Lessons from political philosophy. *Proceedings of the 2018 Conference on Fairness, Accountability, and Transparency*, 149–159. <https://doi.org/10.1145/3287560.3287598>

9. Boutilier, C., Patrascu, R., Poupart, P., & Schuurmans, D. (2006). Constraint-based optimization and utility elicitation in multiattribute decision problems. *IJCAI*, 6, 190–197. <https://doi.org/10.5555/1625275.1625307>
10. Zhao, T., Qin, Z., & Yang, Y. (2021). A survey on explainable artificial intelligence (XAI): Toward transparent and interpretable AI systems. *IEEE Transactions on Neural Networks and Learning Systems*. <https://doi.org/10.1109/TNNLS.2021.3070846>
11. Jurafsky, D., & Martin, J. H. (2023). *Speech and Language Processing* (3rd ed.). Draft. <https://web.stanford.edu/~jurafsky/slp3/>
12. Picard, R. W. (2003). Affective computing: Challenges. *International Journal of Human-Computer Studies*, 59(1–2), 55–64. [https://doi.org/10.1016/S1071-5819\(03\)00052-1](https://doi.org/10.1016/S1071-5819(03)00052-1)
13. Lowe, R., Wu, Y., Tamar, A., Harb, J., Abbeel, P., & Mordatch, I. (2017). Multi-Agent Actor-Critic for Mixed Cooperative-Competitive Environments. In *Advances in Neural Information Processing Systems (NeurIPS)*. <https://arxiv.org/abs/1706.02275>
14. Foerster, J., Farquhar, G., Afouras, T., Nardelli, N., & Whiteson, S. (2018). Counterfactual Multi-Agent Policy Gradients. In *AAAI Conference on Artificial Intelligence*. <https://arxiv.org/abs/1705.08926>
15. Rahwan, I., Cebrian, M., Obradovich, N., Bongard, J., Bonnefon, J.-F., Breazeal, C., & Lazer, D. (2019). Machine Behaviour. *Nature*, 568(7753), 477–486. <https://doi.org/10.1038/s41586-019-1138-y>
16. Faratin, P., Sierra, C., & Jennings, N. R. (1998). Negotiation Decision Functions for Autonomous Agents. *Robotics and Autonomous Systems*, 24(3-4), 159–182. [https://doi.org/10.1016/S0921-8890\(98\)00029-3](https://doi.org/10.1016/S0921-8890(98)00029-3)
17. Sutton, R. S., & Barto, A. G. (2018). *Reinforcement Learning: An Introduction* (2nd ed.). MIT Press. <http://incompleteideas.net/book/the-book-2nd.html>
18. Silver, D., Schrittwieser, J., Simonyan, K., Antonoglou, I., Huang, A., Guez, A., & Hassabis, D. (2017). Mastering the Game of Go Without Human Knowledge. *Nature*, 550(7676), 354–359. <https://doi.org/10.1038/nature24270>
19. Ghosh, S., Basu, S., & Dey, L. (2022). Synthetic Data Generation: Techniques, Challenges, and Applications. *ACM Computing Surveys*. <https://doi.org/10.1145/3502266>
20. Boutilier, C., Patrascu, R., Poupart, P., & Schuurmans, D. (2006). Constraint-Based Optimization and Utility Elicitation in Multiattribute Decision Problems.
21. Lowe, R., Wu, Y., Tamar, A., Harb, J., Abbeel, P., & Mordatch, I. (2017). Multi-agent actor-critic for mixed cooperative-competitive environments. *Advances in Neural Information Processing Systems (NeurIPS)*. <https://doi.org/10.48550/arXiv.1706.02275>
22. Rahwan, I., Cebrian, M., Obradovich, N., et al. (2019). Machine behaviour. *Nature*, 568(7753), 477–486. <https://doi.org/10.1038/s41586-019-1138-y>

23. Faratin, P., Sierra, C., & Jennings, N. R. (1998). Negotiation decision functions for autonomous agents. *Robotics and Autonomous Systems*, 24(3–4), 159–182.
[https://doi.org/10.1016/S0921-8890\(98\)00029-3](https://doi.org/10.1016/S0921-8890(98)00029-3)
24. Foerster, J. N., Farquhar, G., Afouras, T., Nardelli, N., & Whiteson, S. (2018). Counterfactual multi-agent policy gradients. *Proceedings of the AAAI Conference on Artificial Intelligence*, 32(1). <https://doi.org/10.48550/arXiv.1705.08926>
25. Sutton, R. S., & Barto, A. G. (2018). *Reinforcement Learning: An Introduction (2nd ed.)*. MIT Press.
<http://incompleteideas.net/book/the-book-2nd.html>
26. Ghosh, S., Basu, S., & Dey, L. (2022). Synthetic data generation: Techniques, challenges, and applications. *ACM Computing Surveys*, 55(6), 1–40.
<https://doi.org/10.1145/3502266>
27. Boutilier, C., Patrascu, R., Poupart, P., & Schuurmans, D. (2006). Constraint-based optimization and utility elicitation in multi-attribute decision problems. *IJCAI*, 6, 190–197. <https://doi.org/10.5555/1625275.1625307>
28. Rahwan, I., Cebrian, M., Obradovich, N., Bongard, J., Bonnefon, J.-F., reazeal, C., & Lazer, D. (2019). Machine behaviour. *Nature*, 568(7753), 477–486.
<https://doi.org/10.1038/s41586-019-1138-y>
29. Wachter, S., Mittelstadt, B., & Floridi, L. (2017). Why a right to explanation of automated decision-making does not exist in the General Data Protection Regulation. *International Data Privacy Law*, 7(2), 76–99.
<https://doi.org/10.1093/idpl/ix005>
30. Boutilier, C., Patrascu, R., Poupart, P., & Schuurmans, D. (2006). Constraint-based optimization and utility elicitation in multiattribute decision problems. *IJCAI*, 6, 190–197.
<https://doi.org/10.5555/1625275.1625307>
31. Raji, I. D., Binns, R., Venkatasubramanian, S., & Gebru, T. (2020). Saving face: Investigating the ethical concerns of facial recognition auditing. *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency (FAT)**, 145–155.
<https://doi.org/10.1145/3351095.3372874>
32. Organization for Economic Co-operation and Development (OECD). (2021). OECD AI Principles. <https://oecd.ai/en/dashboards/ai-principles>
33. Zavolokina, L., Schlegel, M., & Schwabe, G. (2020). Is fintech a hype? A grounded theory of the ecosystem and strategic implications. *Decision Support Systems*, 133, 113362.
<https://doi.org/10.1016/j.dss.2020.113362>
34. Preukschat, A., & Reed, D. (2021). *Self-Sovereign Identity: Decentralized digital identity and verifiable credentials*. Manning Publications.
35. World Economic Forum. (2021). A blueprint for digital identity: The role of financial institutions in building digital identity. https://www3.weforum.org/docs/WEF_A_Blueprint_for_Digital_Identity.pdf

36. Zwitter, A., & Boisse-Despiaux, M. (2020). Blockchain for the future of digital identity. *Technology in Society*, 62, 101284. <https://doi.org/10.1016/j.techsoc.2020.101284>
37. Calvo, R. A., & D'Mello, S. (2010). Affect detection: An interdisciplinary review of models, methods, and their applications. *IEEE Transactions on Affective Computing*, 1(1), 18–37. <https://doi.org/10.1109/T-AFFC.2010.1>
38. Picard, R. W. (2003). Affective computing: Challenges. *International Journal of Human-Computer Studies*, 59(1–2), 55–64. [https://doi.org/10.1016/S1071-5819\(03\)00052-1](https://doi.org/10.1016/S1071-5819(03)00052-1)
39. Zavolokina, L., Schlegel, M., & Schwabe, G. (2020). Is fintech a hype? A grounded theory of the ecosystem and strategic implications. *Decision Support Systems*, 133, 113362. <https://doi.org/10.1016/j.dss.2020.113362>
40. Preukschat, A., & Reed, D. (2021). *Self-Sovereign Identity: Decentralized digital identity and verifiable credentials*. Manning Publications.
41. World Economic Forum. (2021). A blueprint for digital identity: The role of financial institutions in building digital identity. https://www3.weforum.org/docs/WEF_A_Blueprint_for_Digital_Identity.pdf
42. Zwitter, A., & Boisse-Despiaux, M. (2020). Blockchain for the future of digital identity. *Technology in Society*, 62, 101284. <https://doi.org/10.1016/j.techsoc.2020.101284>
43. Calvo, R. A., & D'Mello, S. (2010). Affect detection: An interdisciplinary review of models, methods, and their applications. *IEEE Transactions on Affective Computing*, 1(1), 18–37. <https://doi.org/10.1109/T-AFFC.2010.1>
44. Picard, R. W. (2003). Affective computing: Challenges. *International Journal of Human-Computer Studies*, 59(1–2), 55–64. [https://doi.org/10.1016/S1071-5819\(03\)00052-1](https://doi.org/10.1016/S1071-5819(03)00052-1)
45. Rahwan, I., Cebrian, M., Obradovich, N., et al. (2019). Machine behaviour. *Nature*, 568(7753), 477–486. <https://doi.org/10.1038/s41586-019-1138-y>
46. Ghosh, S., Basu, S., & Dey, L. (2022). Synthetic data generation: Techniques, challenges, and applications. *ACM Computing Surveys*, 55(6), 1–40. <https://doi.org/10.1145/3502266>
47. Floridi, L., & Cowls, J. (2019). A unified framework of five principles for AI in society. *Harvard Data Science Review*, 1(1). <https://doi.org/10.1162/99608f92.8cd550d1>
48. OECD. (2021). OECD AI Principles. Organisation for Economic Co-operation and Development. <https://oecd.ai/en/dashboards/ai-principles>

Figures




Figure 1

System Architecture




Figure 2

System model diagram




Figure 3

Code Workflow

3D Negotiation Trajectory




Figure 4

Negotiation Trajectory

Consent & Control Architecture




Figure 5

Consent & Control Architecture

Future Work




Figure 6

Future Roadmap